

Electronics and Robotics Club, IIT Bombay

---

# CODE THE PIXELS

An introduction to IMAGE PROCESSING

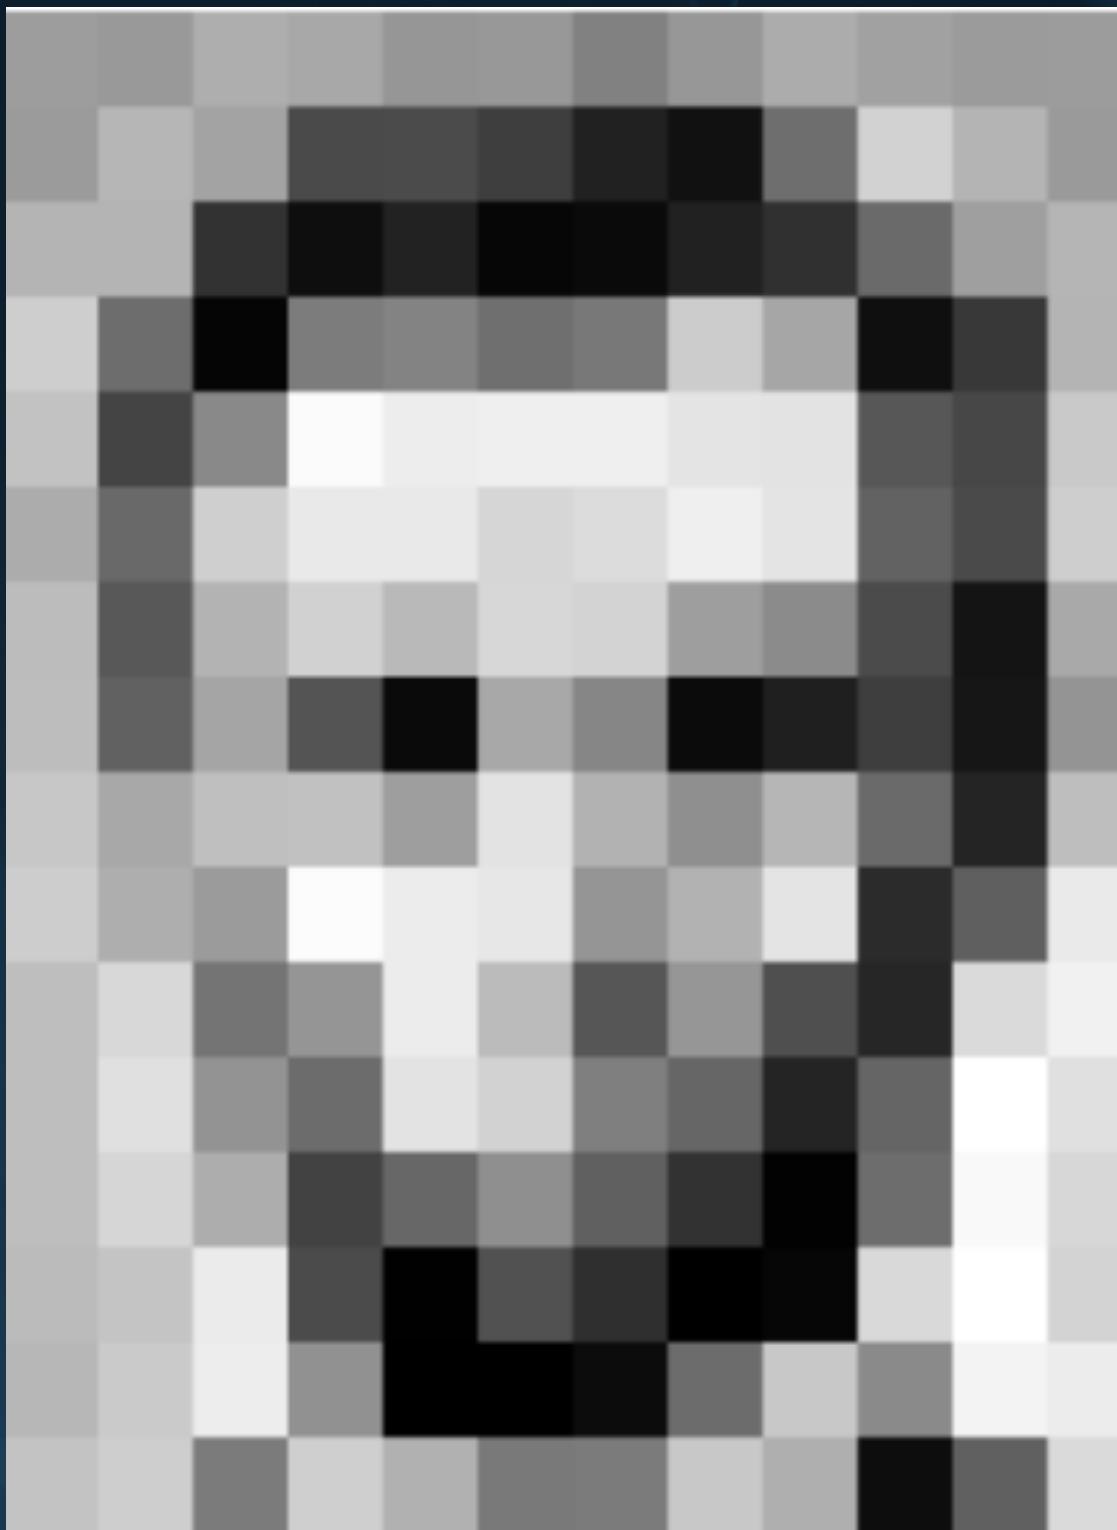
<http://bit.ly/erc-ip>

# OVERVIEW

1. What are Images and What is Image Processing?
2. Reading images and videos in OpenCV
3. Thresholding
4. Colour/Object Detection using OpenCV
5. Blurring
6. Morphology
7. Edge Detection
8. Contours
9. Applications and Uses of Image Processing
10. Going "Deep"

# Images as seen by Computers

- The digital images are not continuous
- They are divided into small blocks known as pixels(short for picture elements)
- Each pixel has a intensity value
  - Range: 0 to 255
  - 8-bit
- More the pixels, more the variations and so better the quality
- (Now you know why megapixel of camera matters)



157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	45	105	159	181
206	109	5	124	191	111	120	204	165	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	35	101	255	224
190	214	173	66	103	143	95	50	2	109	249	218
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	76	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

- More the intensity, whiter the pixel
  - 255 = pure white
  - 0 = pure black

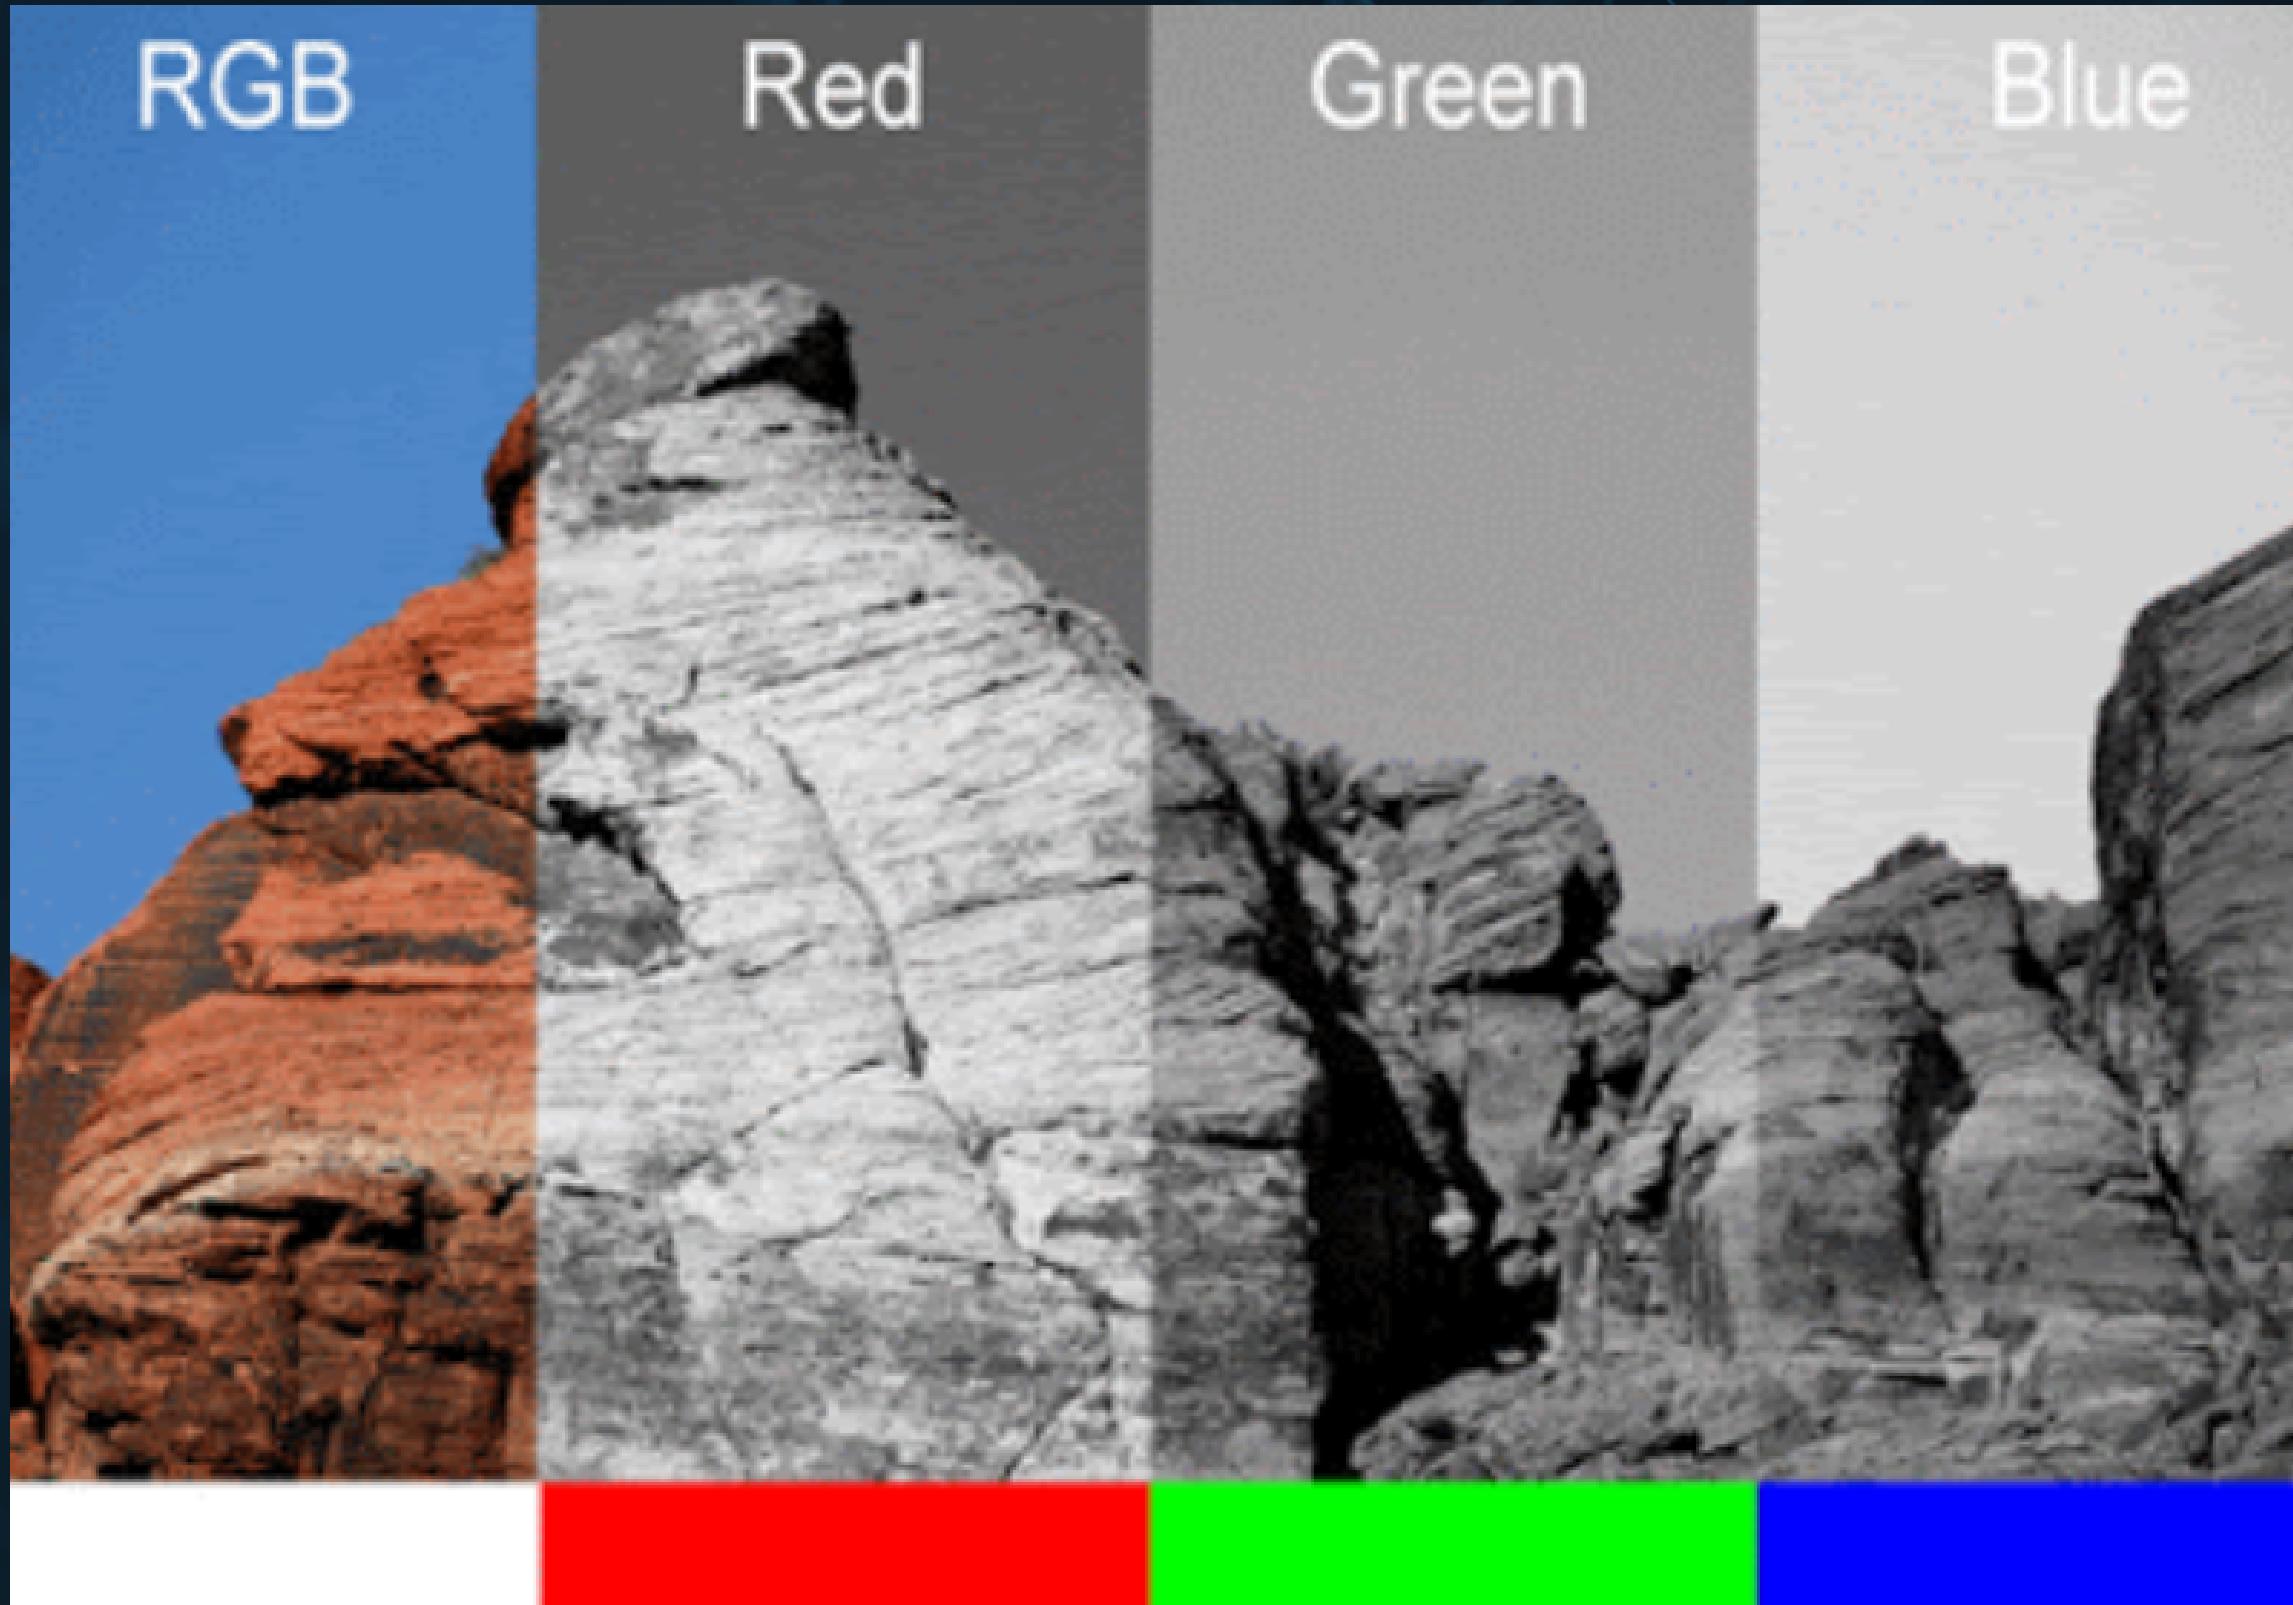
# image processing



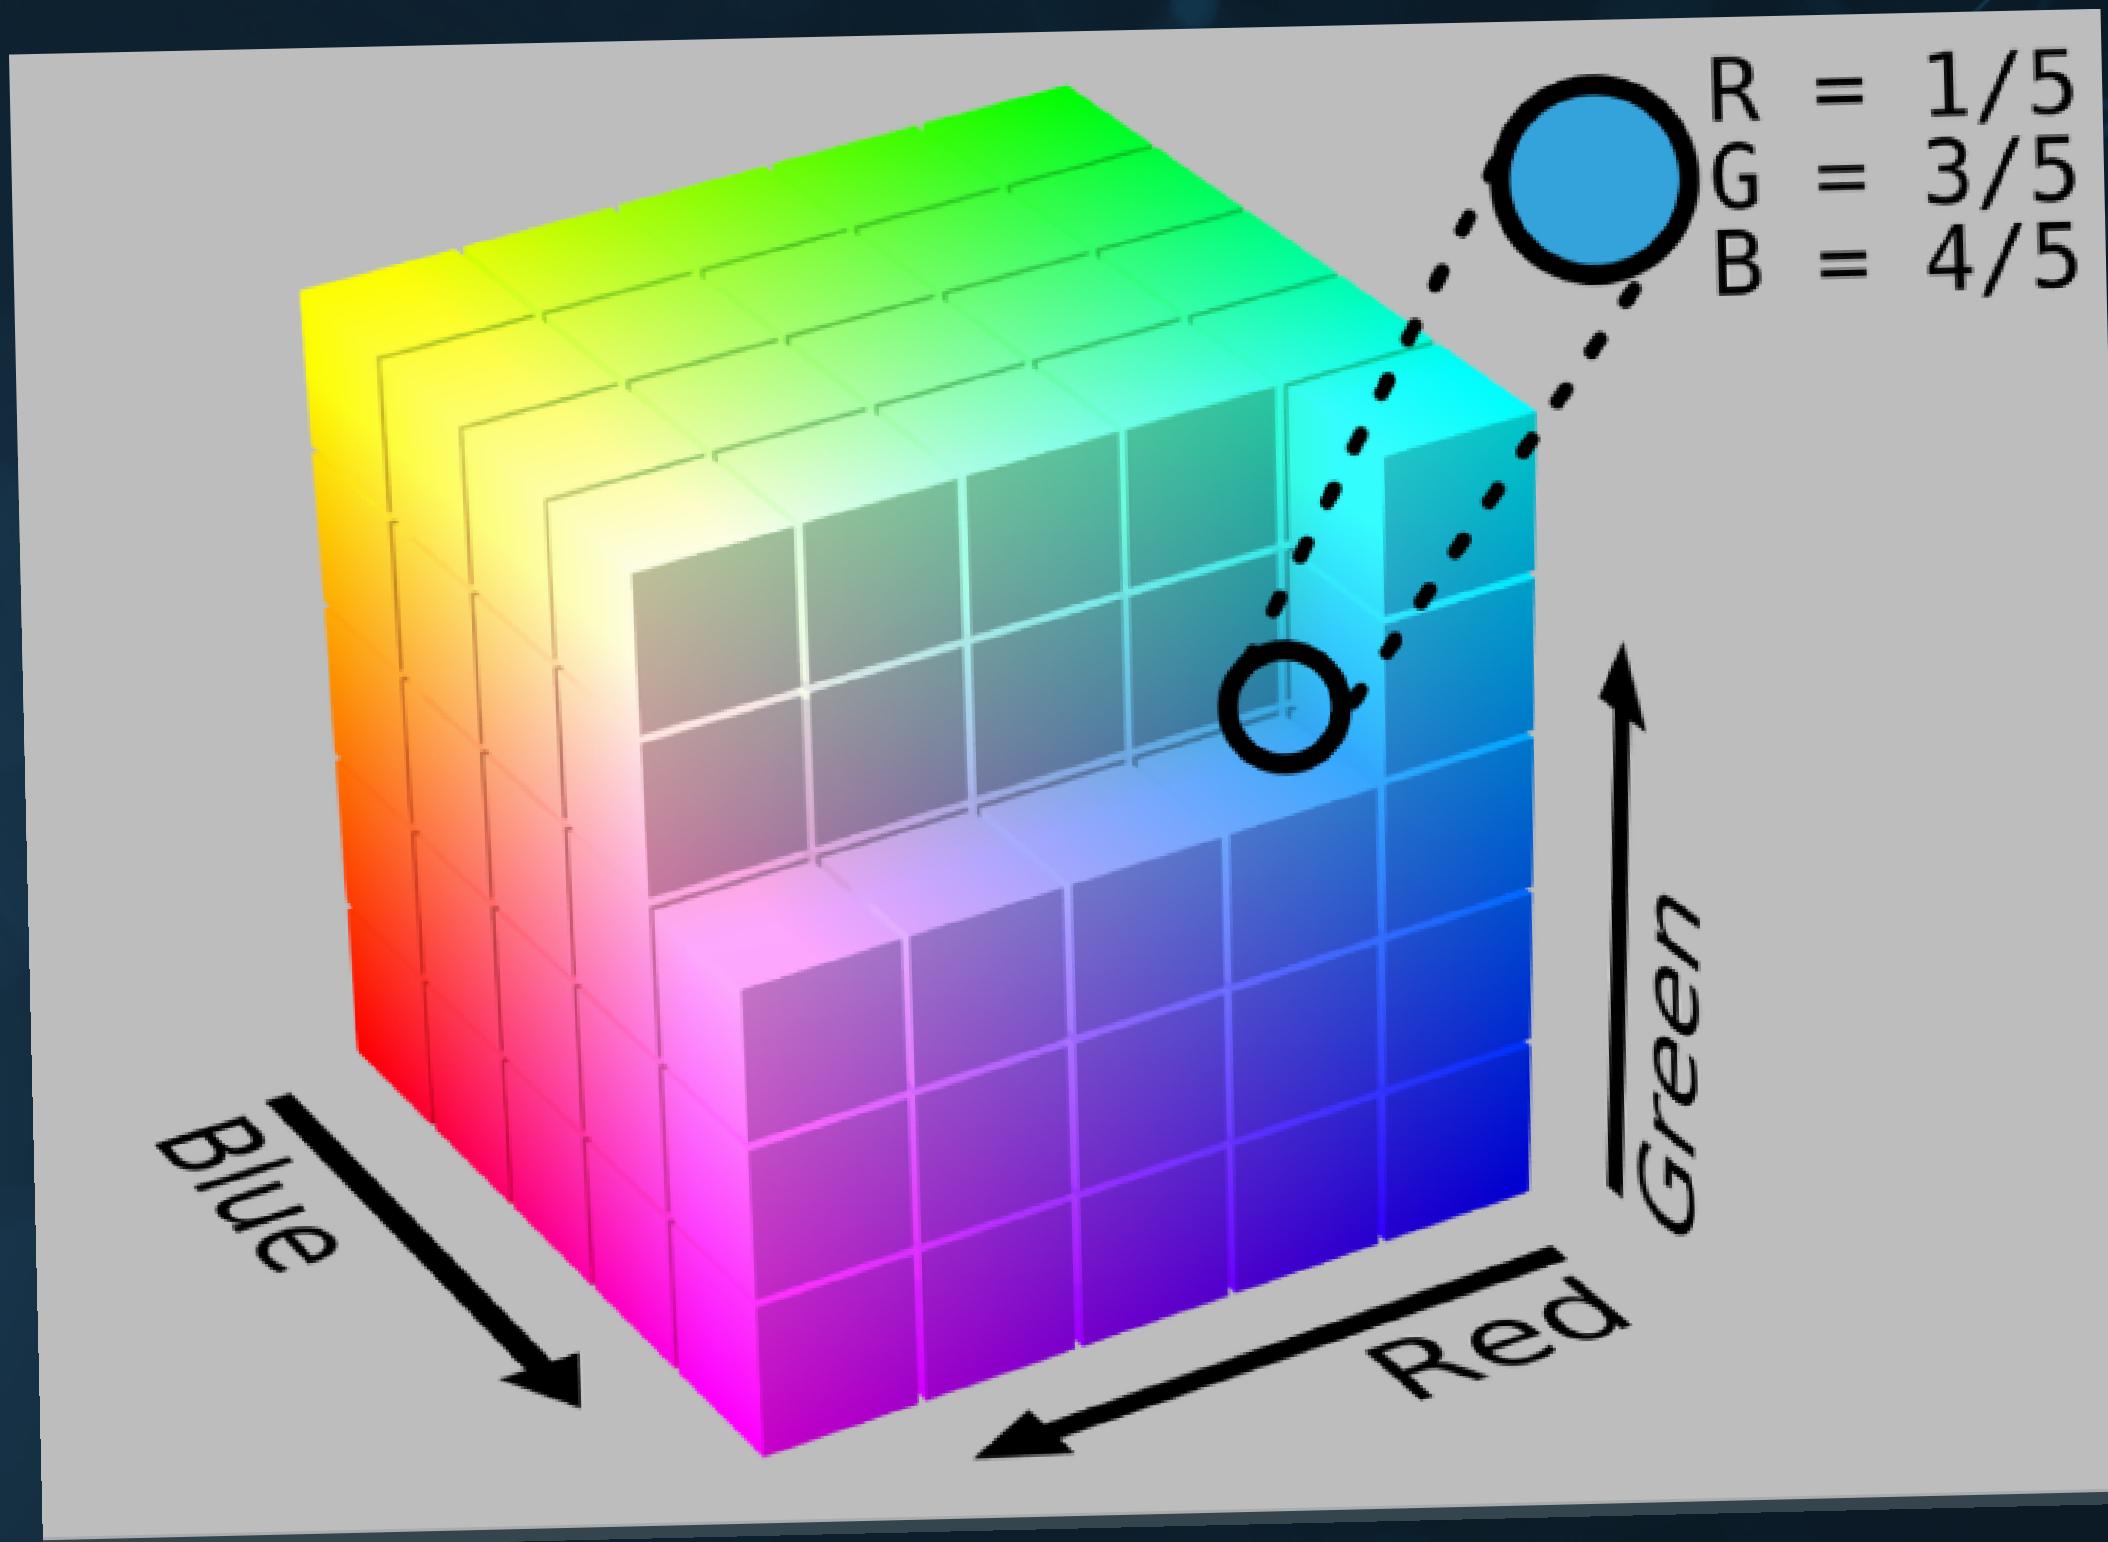
A hand is holding a smartphone, which displays a collage of various technology-related words. The words are arranged in a cloud-like pattern and include: technology, data, internet, digital, graphic, computer, media, art, choice, laptop, success, devices, concepts, work, innovation, advertising, solution, layout, elements, composition, presentation, business, process, options, information, modern, responsibility, trend, employment, critical, instruction, education, search, engineering, reusable, design, website, nature, people, future, creative, environment, professionals, chart, inspiration, world, dramatic, color, exchange, social, step, label, tech.

# Colors!!!

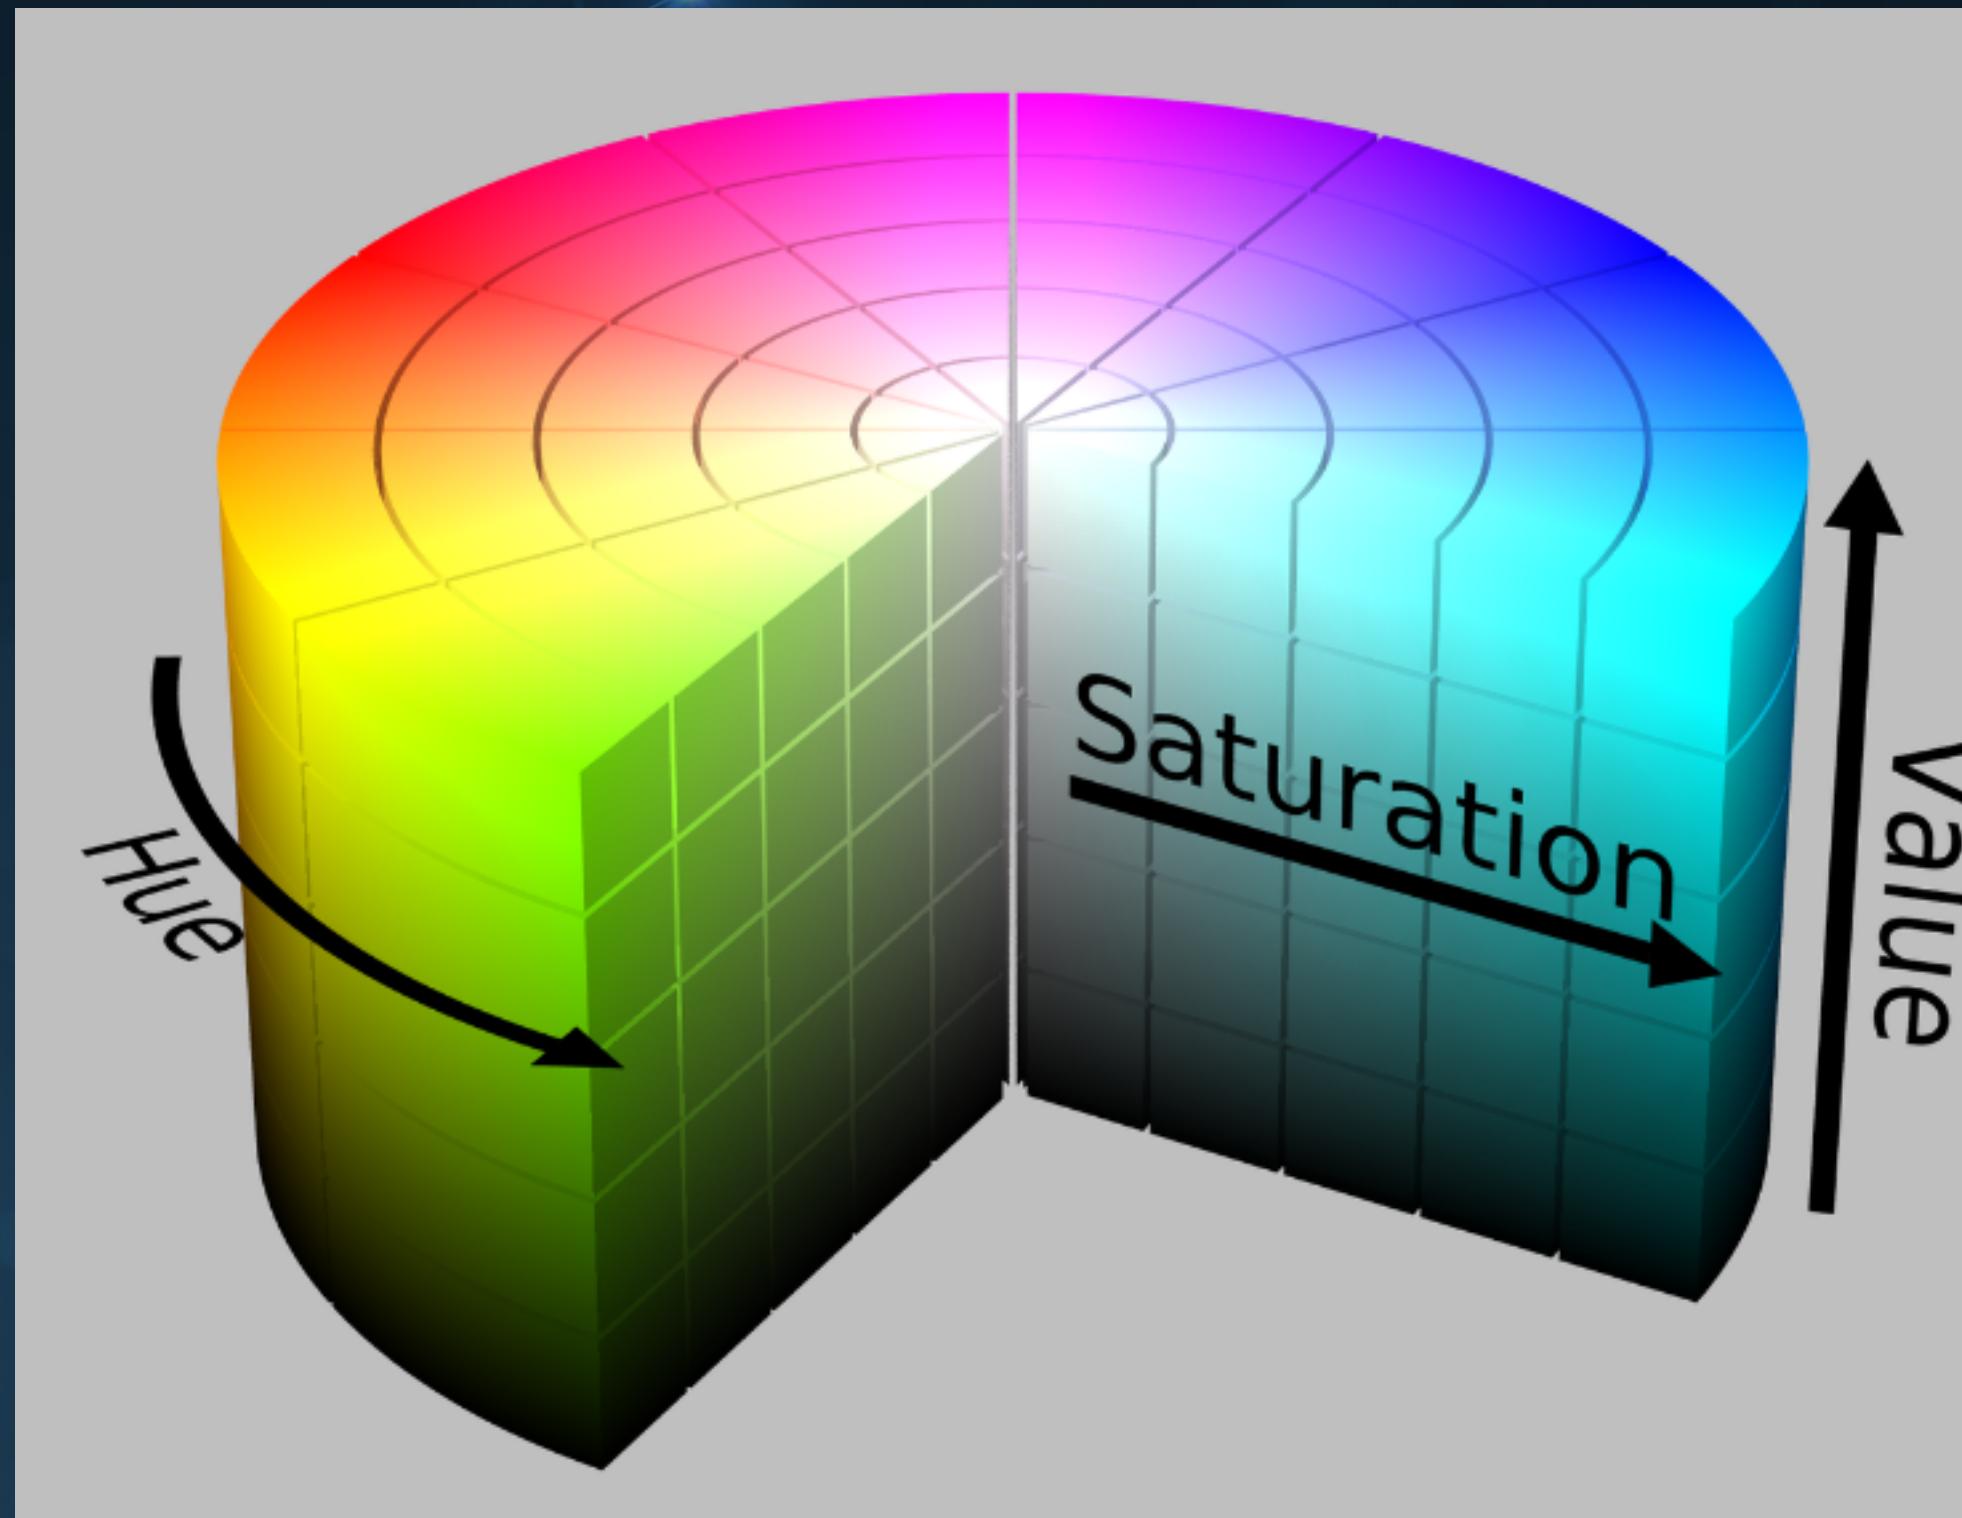
- Pixel no longer has 1 value
- Multiple Channels to accommodate colors
  - Each channel is same as a grayscale image



# Color Formats



RGB format -  
Red, Green, Blue format  
OpenCV uses reverse  
format i.e. BGR.

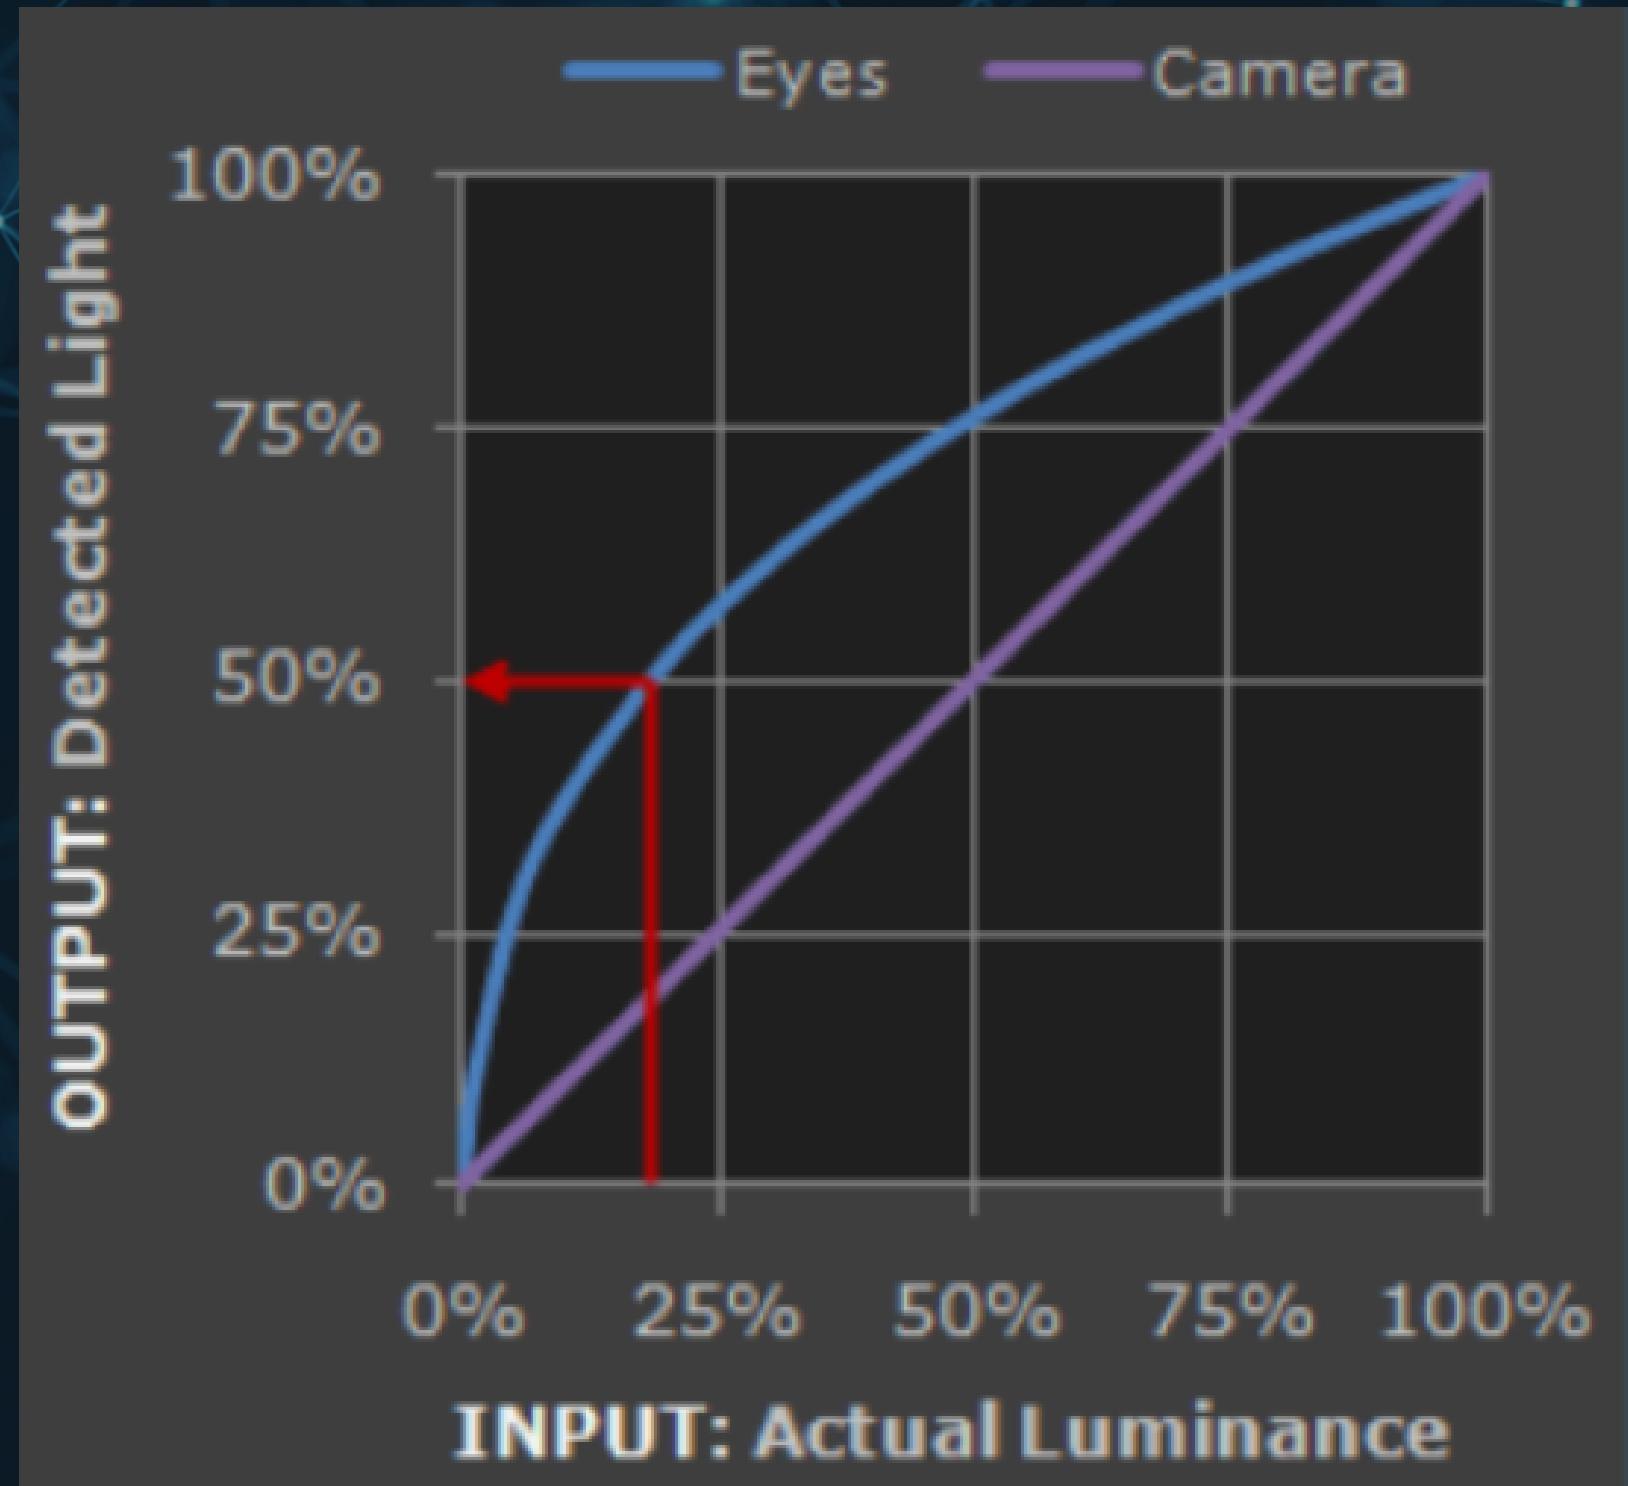


**HSV format -**

Hue, Saturation, Value format

# Gamma Correction

- The human eye response is not linear with respect to intensity
- It responds more to lower intensity level changes
- Power Law:  
$$\text{Detected} = (\text{Luminance})^{\gamma}$$
- Correction is done by taking the root





$\gamma = 1/3$

$\gamma = 2$

# OpenCV Basics

# Reading an Image in OpenCV

```
import cv2  
  
#Reading an image  
image = cv2.imread("joker_1.jpg", 1)  
  
#Resizing the image because it is huge!  
image = cv2.resize(image, (720, 720))  
  
#Displaying the image  
cv2.imshow('window', image)  
  
k = cv2.waitKey(0) #imshow does not render without waitKey  
if k == 95:  
    cv2.destroyAllWindows()
```

cv2.imread(file\_name, flag)

FLAG:

- 0 - Grayscale
- 1 - Coloured(BGR) image
- 1 - Coloured image with alpha channels

# Showing an Image in OpenCV

```
import cv2

#Reading an image
image = cv2.imread("joker_1.jpg", 1)

#Resizing the image because it is huge!
image = cv2.resize(image, (720, 720))

#Displaying the image
cv2.imshow('window', image)

k = cv2.waitKey(0) #imshow does not render without waitKey
if k == 95:
    cv2.destroyAllWindows()
```

cv2.imshow(Window\_name, Image)

Window\_name:  
Name displayed on the display

Image:  
The image to be showed

# Showing an Image in OpenCV

```
import cv2

#Reading an image
image = cv2.imread("joker_1.jpg", 1)

#Resizing the image because it is huge!
image = cv2.resize(image, (720, 720))

#Displaying the image
cv2.imshow('window', image)

k = cv2.waitKey(0) #imshow does not render without waitKey
if k == 95:
    cv2.destroyAllWindows()
```

cv2.waitKey(delay)

delay:

How much time(msec.) to wait before closing the window

0: don't close automatically

returns:

Returns the ASCII value of the key pressed.

# Changing between Color Spaces

```
#Reading a image
frame = cv2.imread("joker.jpg", 1)
#Resizing the image because it is huge!
frame = cv2.resize(frame, (500, 500))

#Converting to HSV space
hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)

cv2.imshow("Showing image", mask)
cv2.imshow("Showing hsv", hsv)
k = cv2.waitKey(0)
if k == 37:
    cv2.destroyAllWindows()
```

cv2.cvtColor(src, code)

src:

Source image

code:

Code to specify the change in color space

General Format:

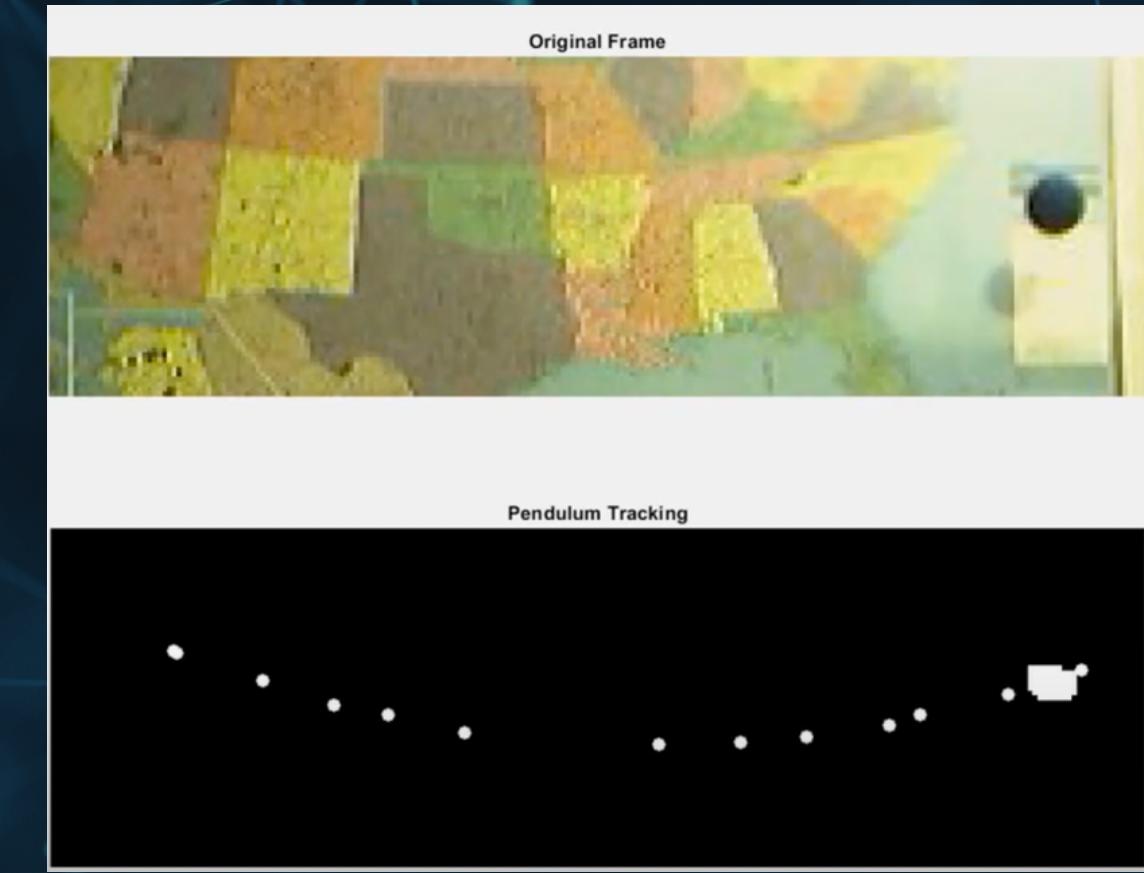
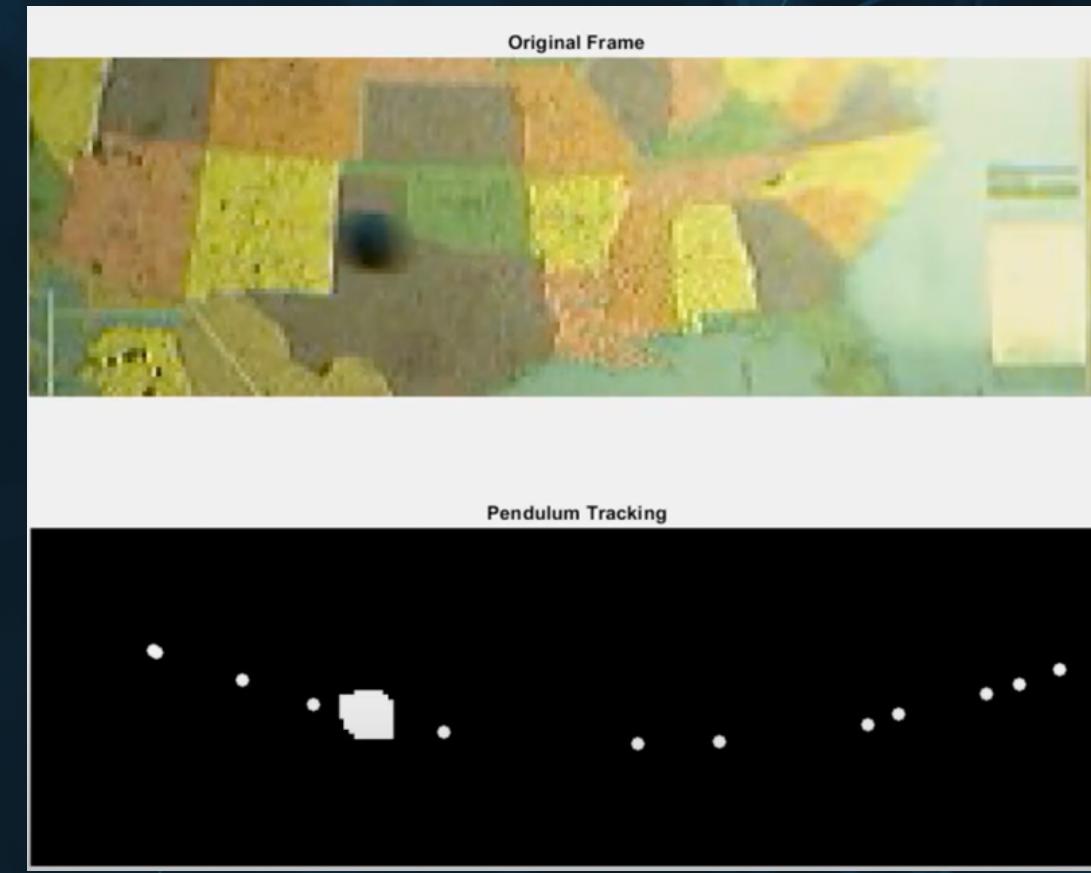
COLOR\_<current>2<wanted>

[https://docs.opencv.org/3.4/d8/d01/group\\_imgproc\\_color\\_conversions.html](https://docs.opencv.org/3.4/d8/d01/group_imgproc_color_conversions.html)

# Real Processing Starts Now!

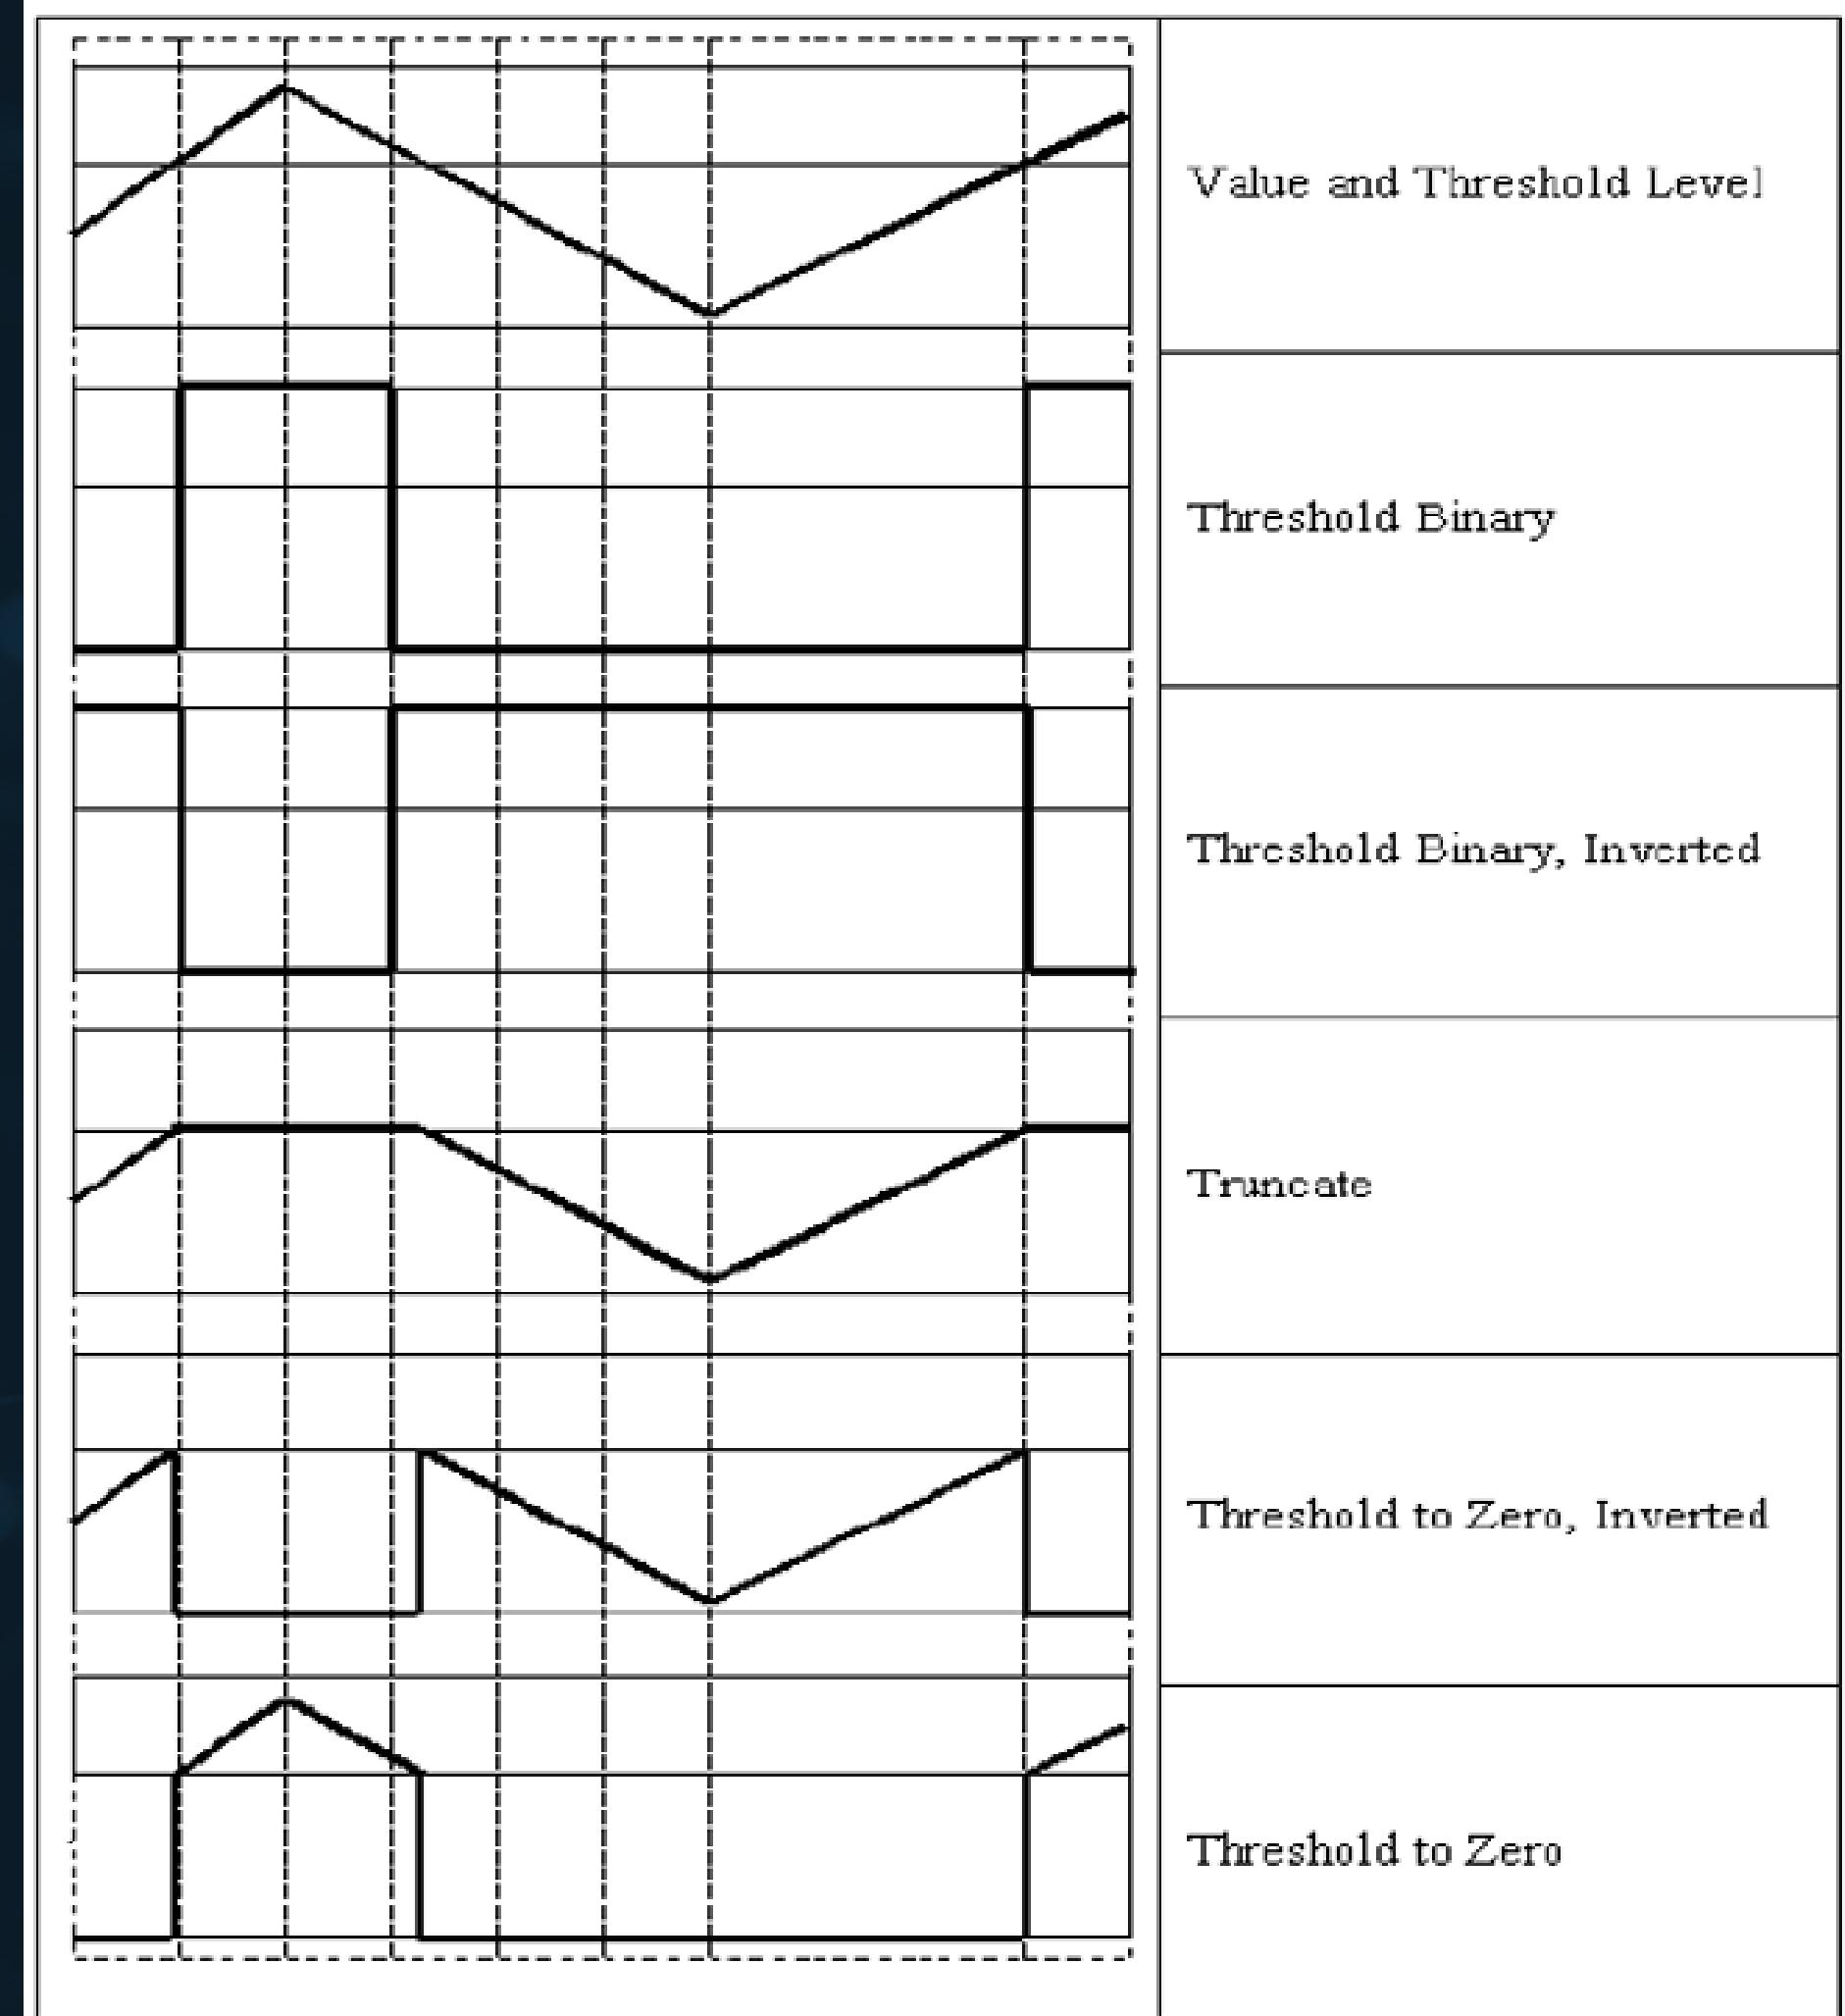
# Thresholding

# Pendulum Tracking using Thresholding



- **Simple Thresholding**
- **Adaptive Thresholding**

# 5 types of thresholding supported by openCV



# Simple Thresholding

- Replace each pixel with some value depending on a given threshold.
- Converts the grayscale image to binary image.

THRESH\_BINARY

Python: cv.THRESH\_BINARY

$$dst(x, y) = \begin{cases} \text{maxval} & \text{if } \text{src}(x, y) > \text{thresh} \\ 0 & \text{otherwise} \end{cases}$$

THRESH\_BINARY\_INV

Python: cv.THRESH\_BINARY\_INV

$$dst(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{maxval} & \text{otherwise} \end{cases}$$

THRESH\_TRUNC

Python: cv.THRESH\_TRUNC

$$dst(x, y) = \begin{cases} \text{threshold} & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$$

THRESH\_TOZERO

Python: cv.THRESH\_TOZERO

$$dst(x, y) = \begin{cases} \text{src}(x, y) & \text{if } \text{src}(x, y) > \text{thresh} \\ 0 & \text{otherwise} \end{cases}$$

THRESH\_TOZERO\_INV

Python: cv.THRESH\_TOZERO\_INV

$$dst(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$$

# Simple Thresholding - OpenCV

`cv2.threshold(src, thresh, maxval, type)`

**src:**

The (grayscale) image to be thresholded

**thresh:**

The threshold to be used

**maxval:**

The maximum value to be used in THRESH\_BINARY

**type:**

Type of thresholding to be used

```
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
ret, thresh1 = cv2.threshold(gray, 127, 255, cv2.THRESH_BINARY)
ret, thresh2 = cv2.threshold(gray, 127, 255, cv2.THRESH_BINARY_INV)
ret, thresh3 = cv2.threshold(gray, 127, 255, cv2.THRESH_TRUNC)
ret, thresh4 = cv2.threshold(gray, 127, 255, cv2.THRESH_TOZERO)
ret, thresh5 = cv2.threshold(gray, 127, 255, cv2.THRESH_TOZERO_INV)
```

# Adaptive Thresholding

- Adaptive thresholding chooses a threshold for a pixel based on its neighboring pixels
- There are two ways to calculate this threshold:
  - Using mean: Given a window size (let's say  $5 \times 5$ ) we calculated the mean for all pixels in that window and then use that as a threshold. Note that you will need to specify what type of thresholding you need (one of the 5 types discussed before)
  - Using weighted mean: This is similar as above only that weights are chosen from a normal distribution

Src maxValue adaptiveMethod thresholdType(only binary ones allowed) blockSize C

#Adaptive threshold

```
adaptive_thresh1 = cv2.adaptiveThreshold(gray, 255, cv2.ADAPTIVE_THRESH_MEAN_C, cv2.THRESH_BINARY, 11, 2)
adaptive_thresh2 = cv2.adaptiveThreshold(gray, 255, cv2.ADAPTIVE_THRESH_GAUSSIAN_C, cv2.THRESH_BINARY, 11, 2)
```

## Image

### Region-based segmentation

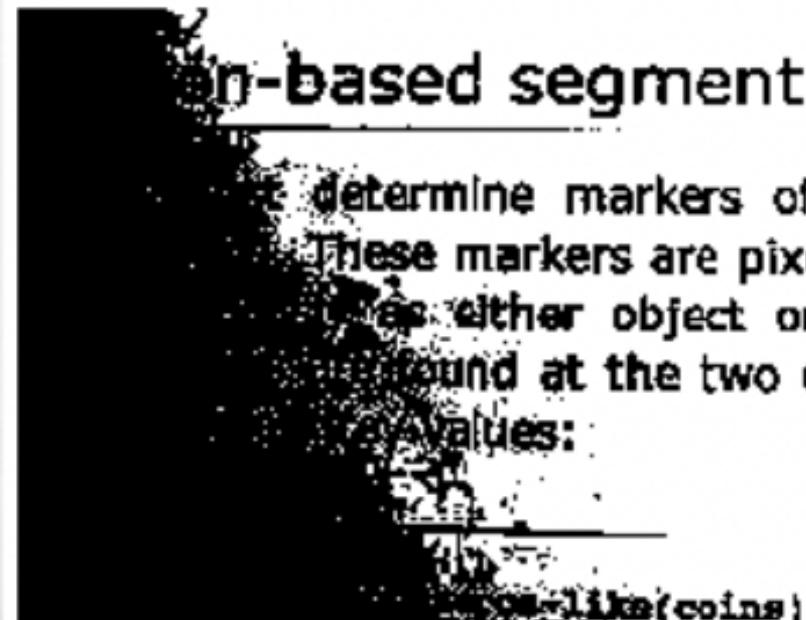
Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

### Global thresholding

#### Region-based segmentation

Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:



### Adaptive thresholding

### Region-based segmentation

Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

**Figure 1:** Top: Original input image. Middle: Applying global thresholding leads to a poor segmentation result. Bottom: Using adaptive thresholding creates a much cleaner segmentation ([image source](#)).

# Color Detection

# Masking



# Masking

- Mask allows us to focus on the portions of the image that interests us.
- Use HSV format and apply filter on Hue channel
- Good Range:  $[h-10, h+10]$
- Saturation and Value are generally kept for whole range. So, change according to need

# Color Detection

**inRange(src, lowerb, upperb)**

**src:**

The image to be thresholded

**lowerb:**

Lower bounds on the pixel intensity values

**upperb:**

Upper bounds on the pixel intensity values

**RETURNS:**

Binary mask

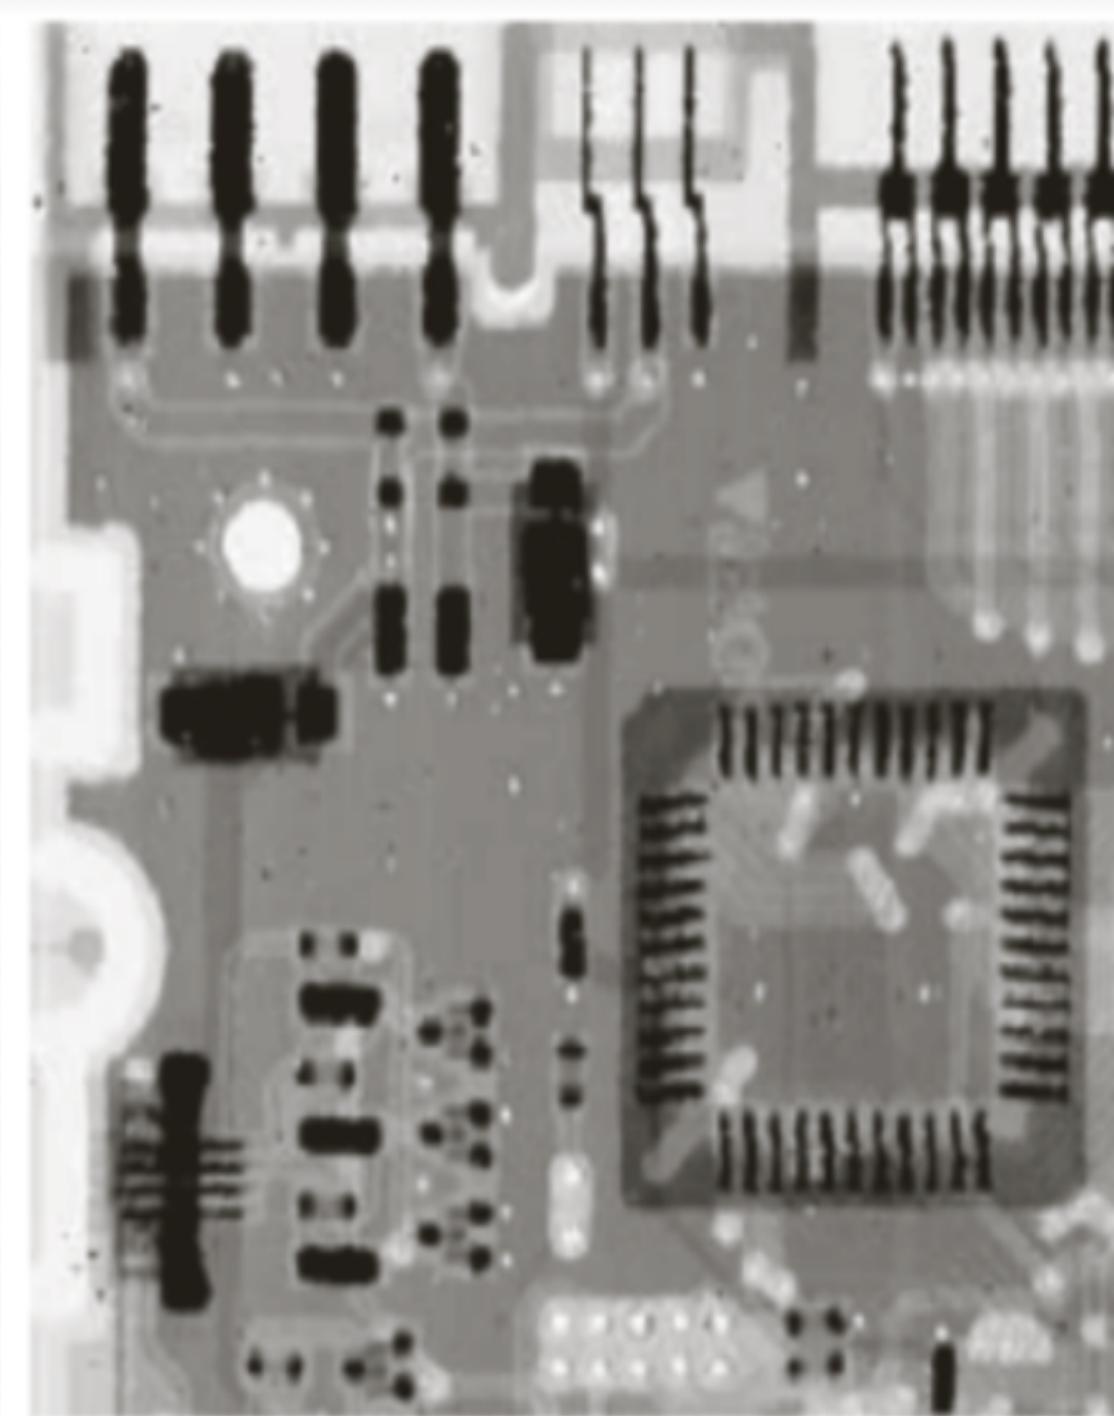
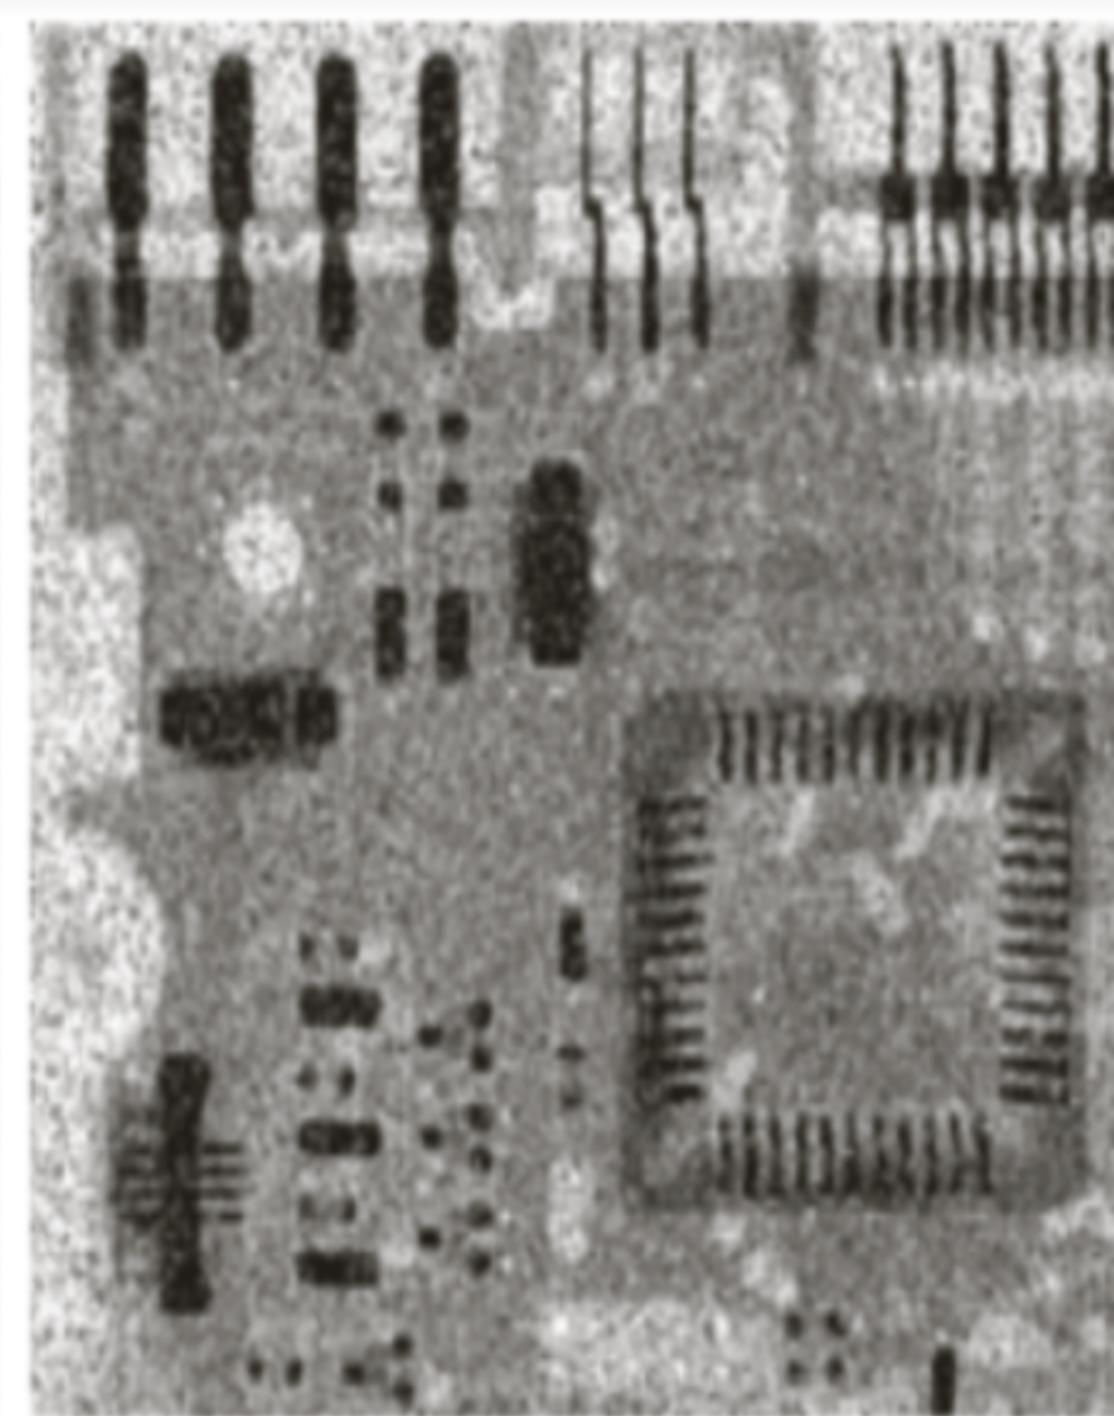
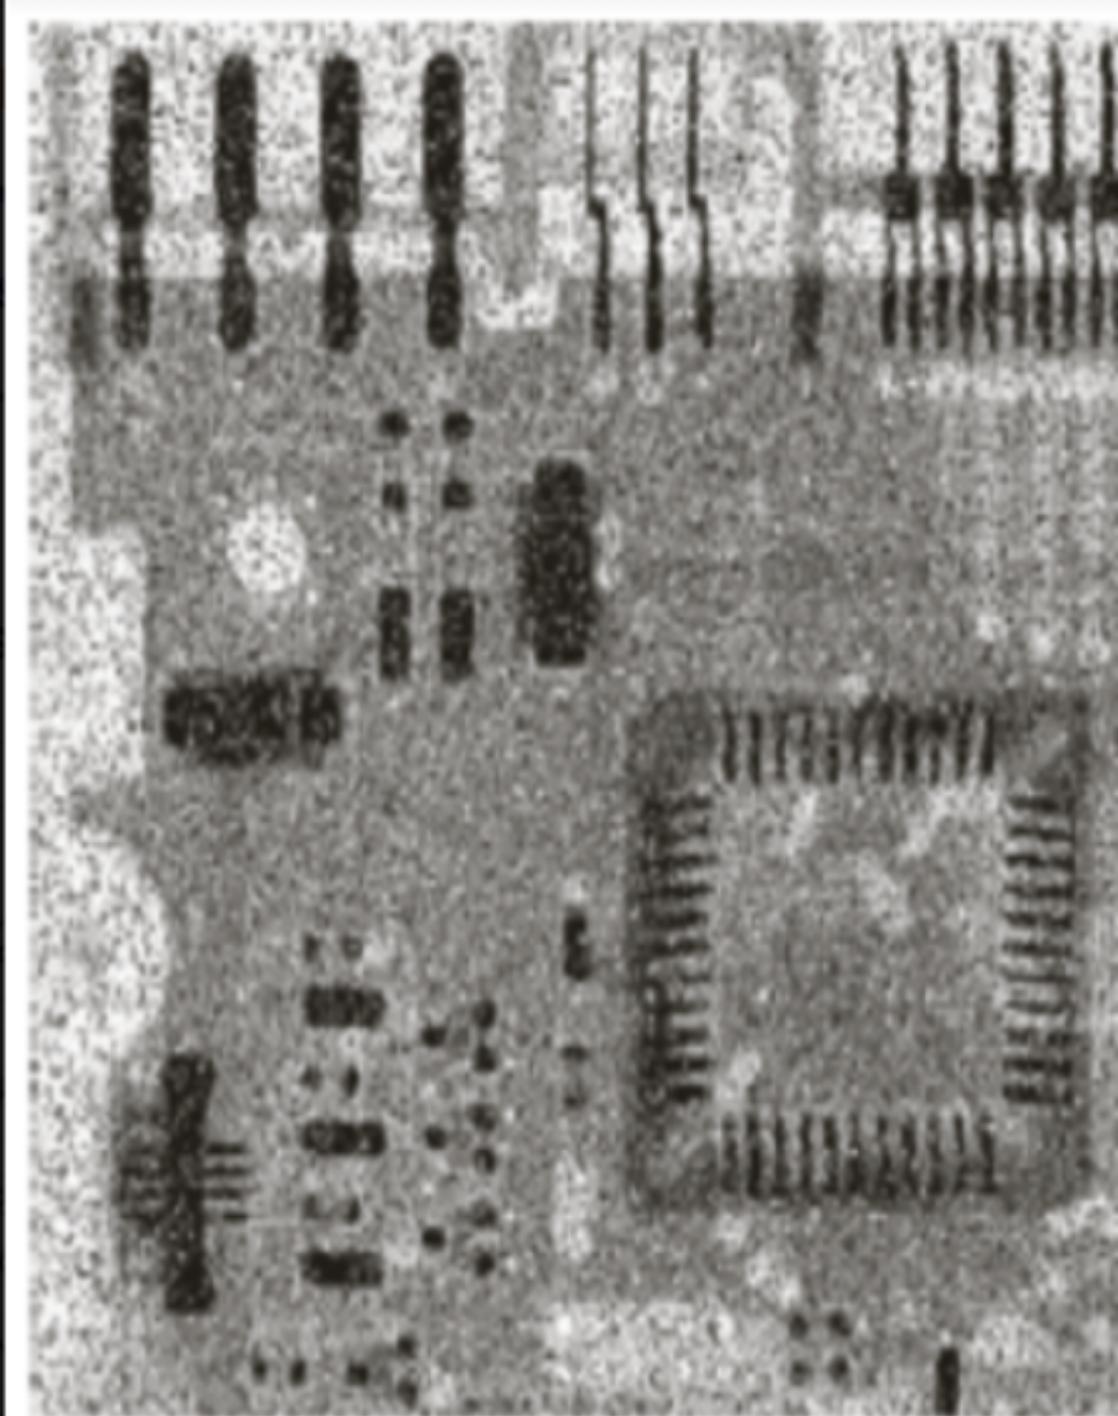
```
l_b = np.array([l_h, l_s, l_v])
u_b = np.array([u_h, u_s, u_v])

mask = cv2.inRange(hsv, l_b, u_b)

res = cv2.bitwise_and(frame, frame, mask=mask)
```

Note: In OpenCV, Hue is between 0 and 180.

# Blurring



a b c

**FIGURE 3.35** (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a  $3 \times 3$  averaging mask. (c) Noise reduction with a  $3 \times 3$  median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)

- Averaging Blur(Box Filter)
- Gaussian Blur(Gaussian Filter)
- Median Blur(Median Filter)

# Kernel and Convolution

- A kernel is nothing but a small matrix, also called convolution matrix.
- Convolution can be thought of as weighted averaging around a pixel
  - Weights = kernel

The diagram illustrates the convolution operation  $I * K$ . It shows three matrices: the input matrix  $I$ , the kernel matrix  $K$ , and the resulting output matrix  $I * K$ .

The input matrix  $I$  is a 7x7 grid:

0	1	1	1	0	0	0	0
0	0	1	1	1	0	0	0
0	0	0	1	1	1	0	0
0	0	0	1	1	0	0	0
0	0	1	1	0	0	0	0
0	1	1	0	0	0	0	0
1	1	0	0	0	0	0	0

The kernel matrix  $K$  is a 3x3 grid:

1	0	1
0	1	0
1	0	1

The resulting output matrix  $I * K$  is a 5x5 grid:

1	4	3	4	1
1	2	4	3	3
1	2	3	4	1
1	3	3	1	1
3	3	1	1	0

Dotted arrows show the receptive field of each output unit in  $I * K$  based on the kernel  $K$ . The result is labeled  $=$ .

$1 \times 1 = t$   
Calculus experts: Ok lol  
Everyone else:



# Kernel and Convolution - Padding

# Averaging Blur(Box Filter)

- The center of a  $n \times n$  kernel is replaced by the average of the  $n^2$  pixels in it.
- The kernel size is always an odd integer.

$$K = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

## blur(image,size)

**image:** Image to be blurred

**size:** height and width of the kernel

**RETURNS:** blurred(convolved) image

```
blurred_image=cv2.blur(image, (3,3))
```

# Gaussian Blur(Gaussian Filter)

- The center of a  $n \times n$  kernel is replaced by the average of the  $n^2$  pixels in it.
- The kernel size is always an odd integer.

$$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

**GaussianBlur(image,size,sigmaX,sigmaY=None,borderType=None)**

**image:** Image to be blurred

**size:** height and width of the kernel

**sigmaX:** standard deviation of the kernel in X direction(if 0, calculated from size)

# Median Blur(Median Filter)

- The center of a  $n \times n$  kernel is replaced by the median of the  $n^2$  pixels in it.
- Highly useful in reducing noise from images.
  - In the above filters, a pixel is set to a calculated value that might not belong to any surrounding pixels. In this method, we set a pixel to value which is present in its neighbourhood. This reduces noise very effectively.

**medianBlur(image,size)**

**image:** Image to be blurred

**size:** size of the window in which median is to be taken

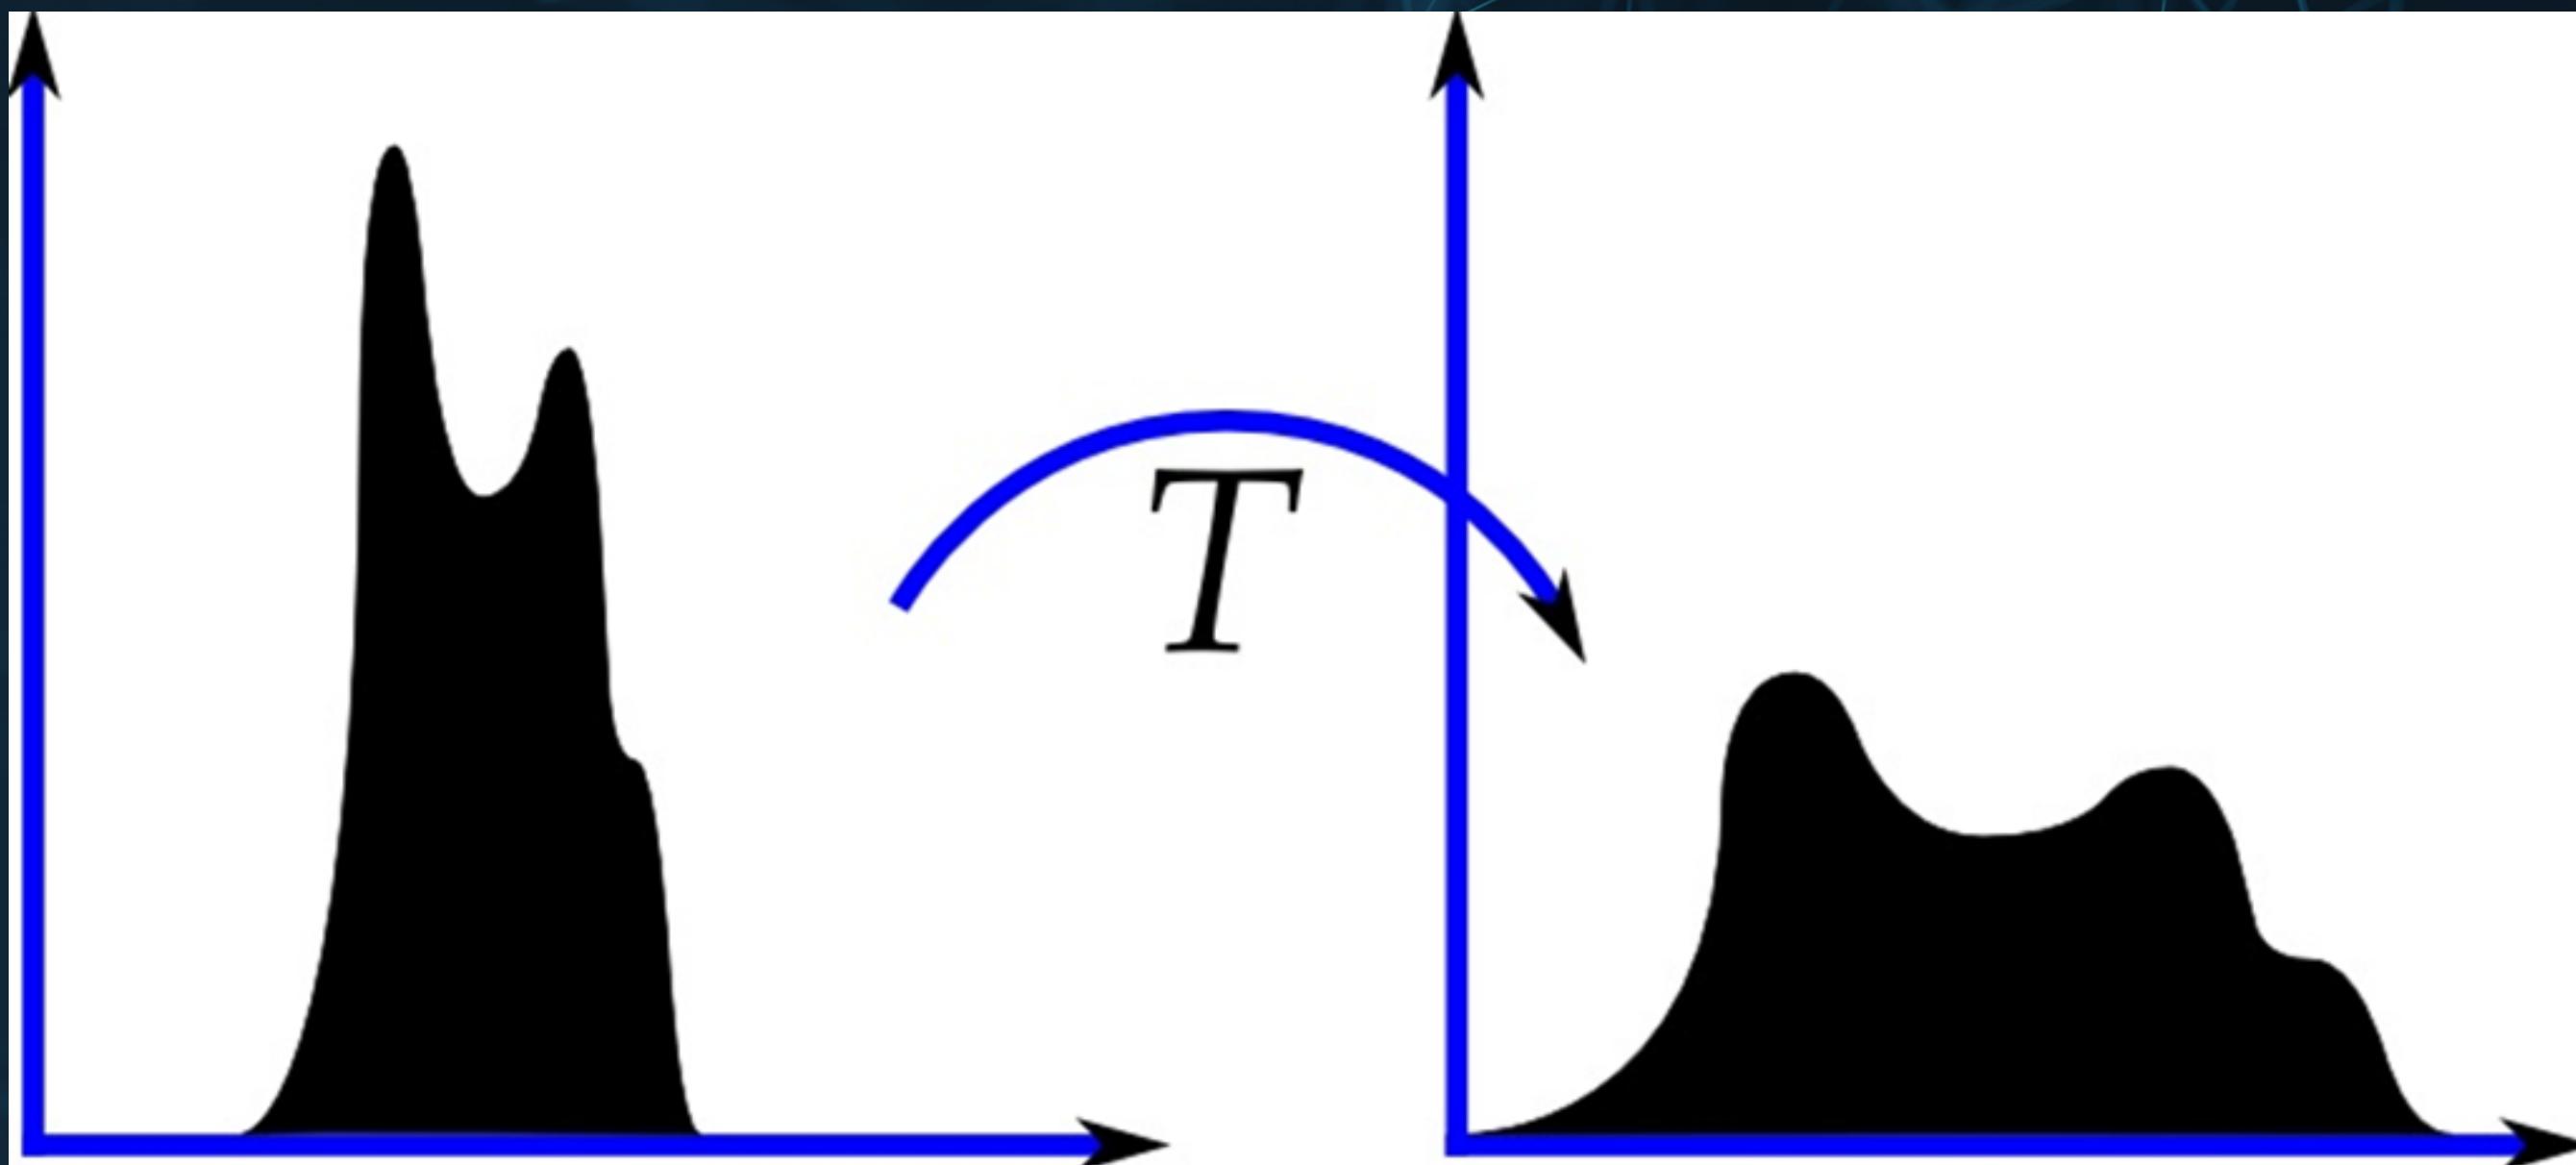
**RETURNS:** blurred(convolved) image

# Noise Reduction using blur

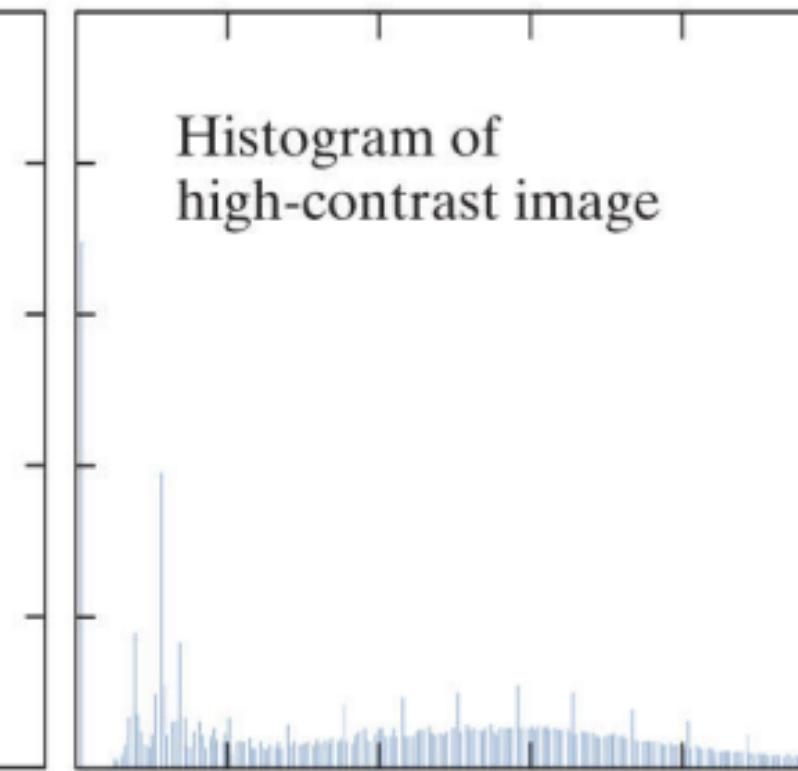
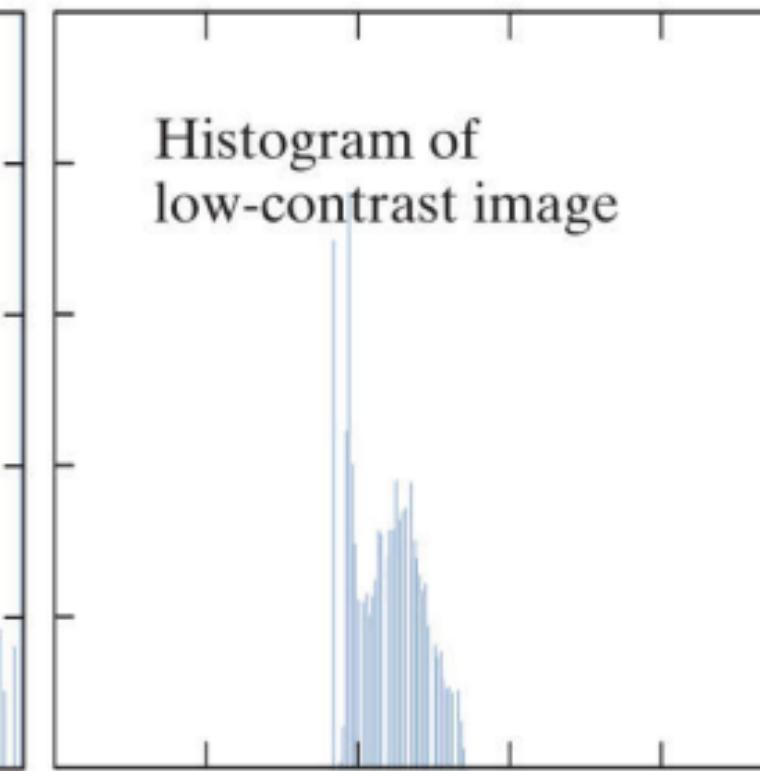
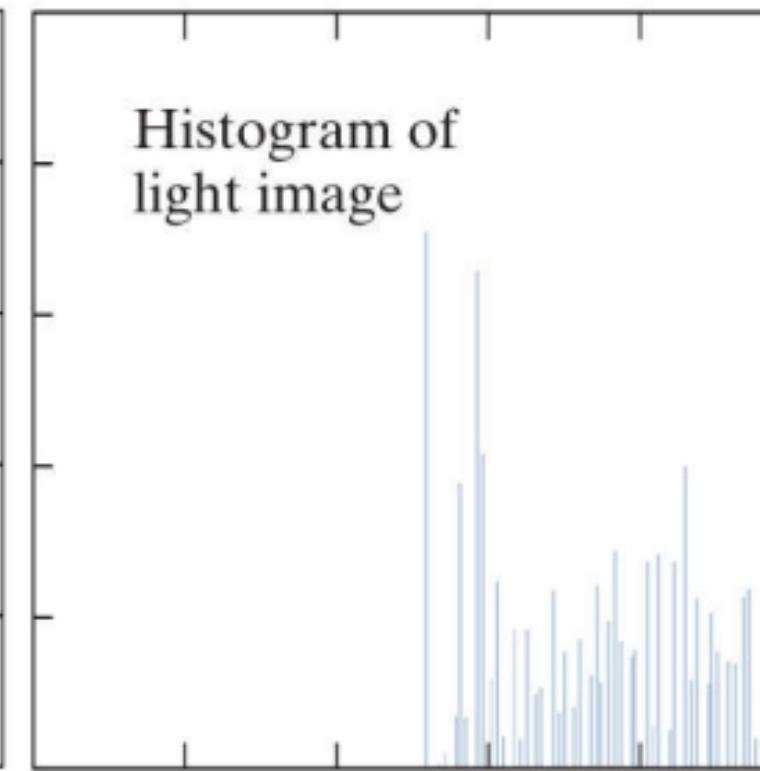
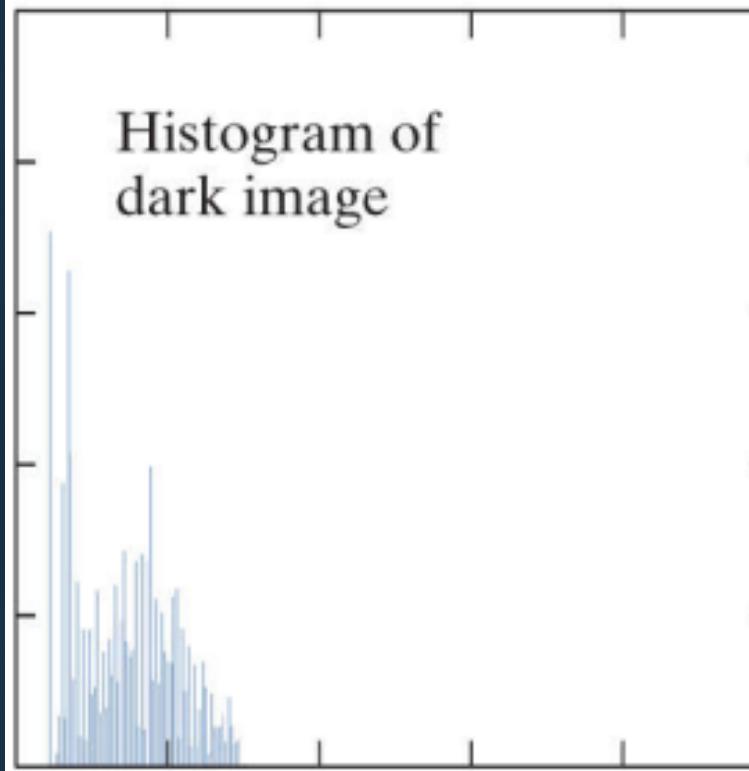
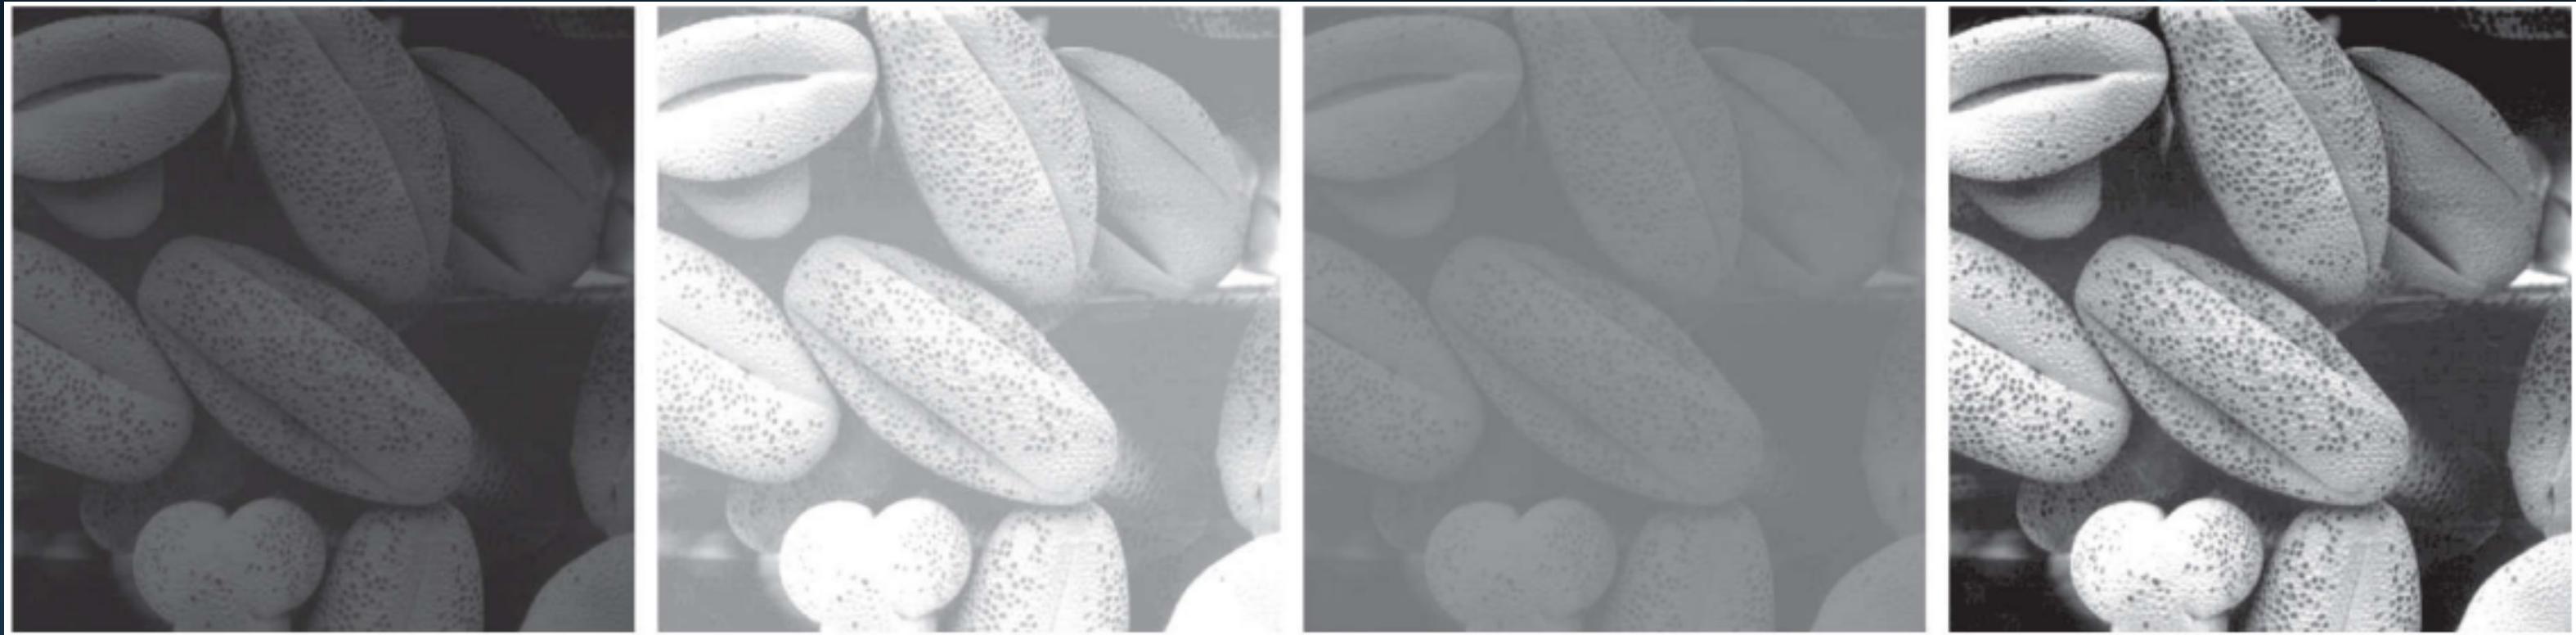
# Some other Processing

# Histogram

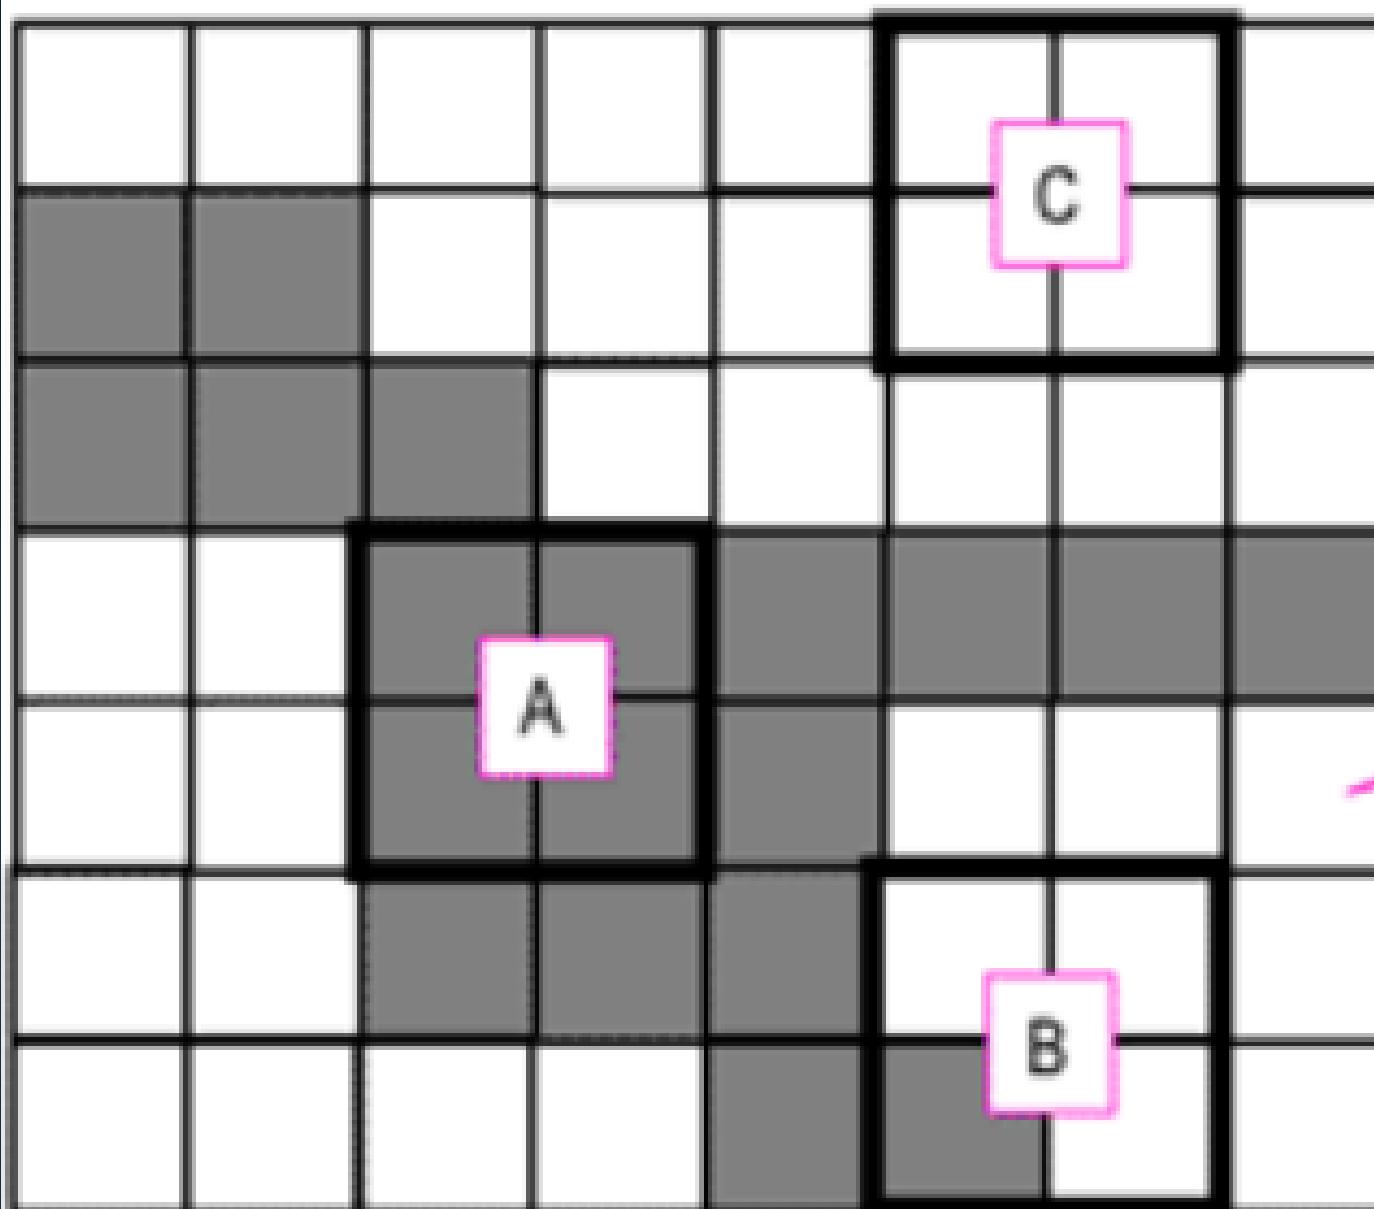
- Number of Pixels vs. Pixel Values
- More variance implies better "contrast"



# Histogram Equalization

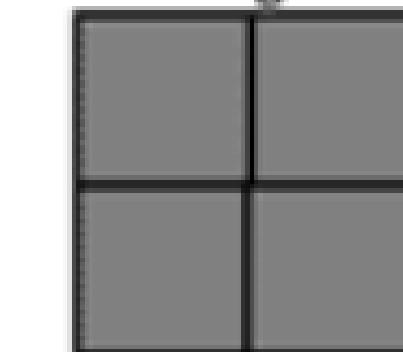


# Morphology



- A - the structuring element fits the image
- B - the structuring element hits (intersects) the image
- C - the structuring element neither fits, nor hits the image

Structuring element



Probing of an image with a structuring element  
(white and grey pixels have zero and non-zero values, respectively).

# Morphology

Erode -->



<-- Dilate

# Morphology

Open - Erode and then Dilate  
Essentially removes white dots

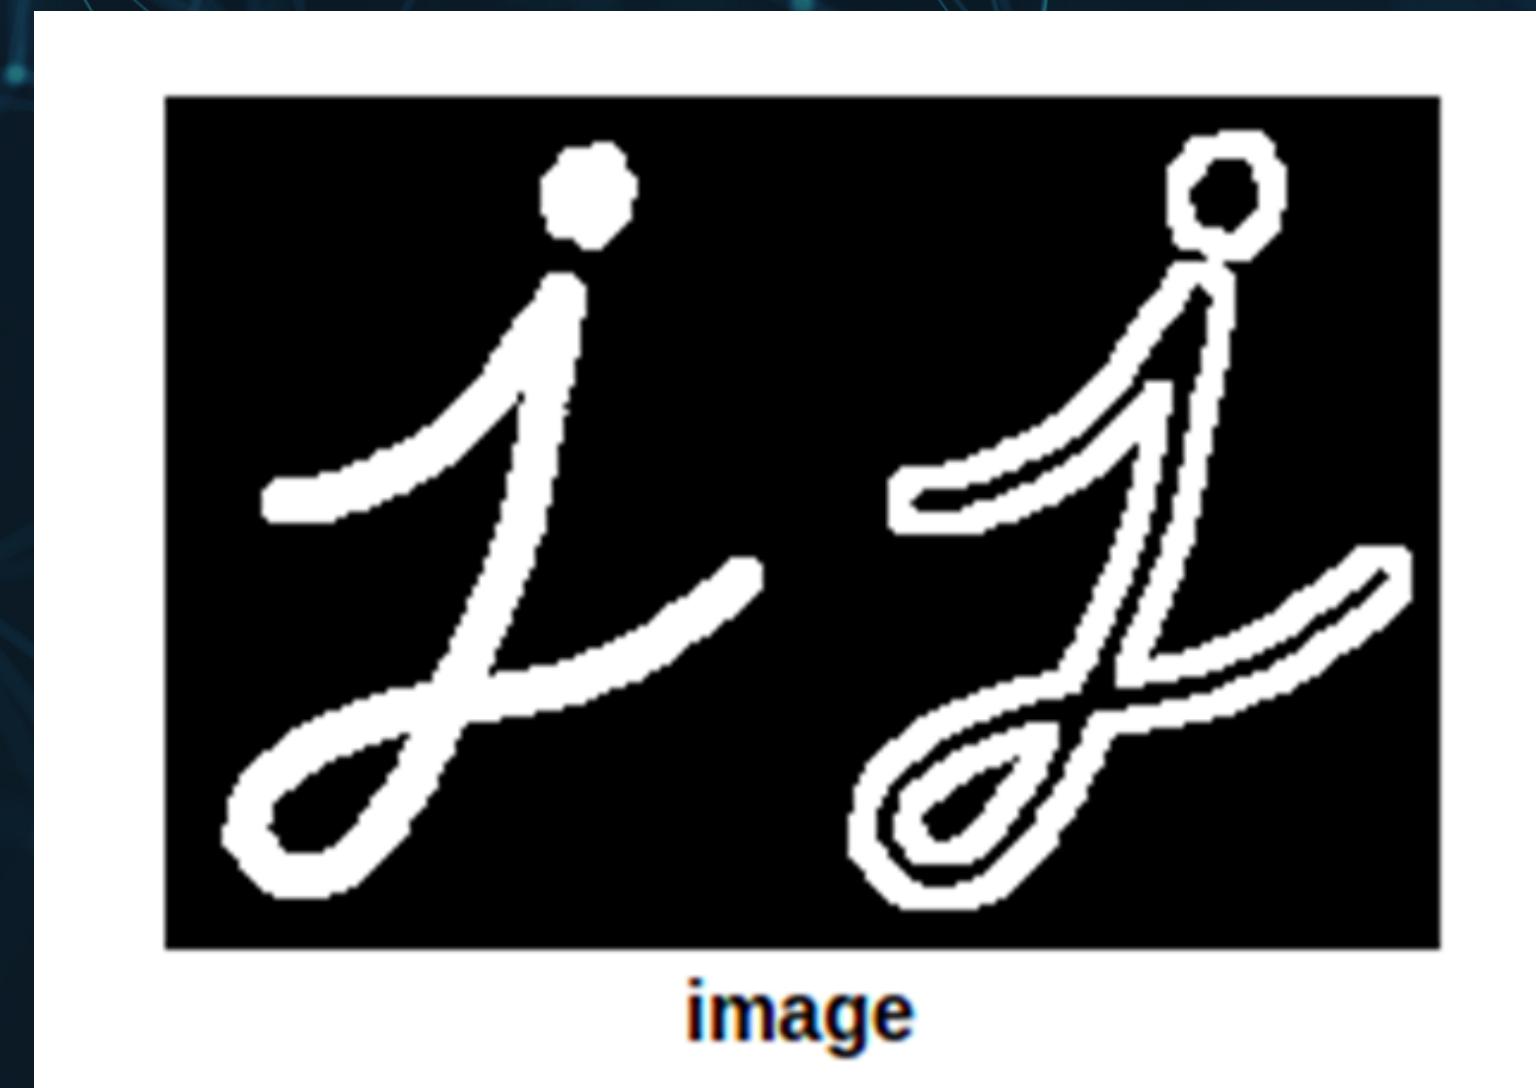


Close - Dilate and then Erode  
Essentially fills up black dots



# Morphology

Gradient - Difference in erosion and dilation



# Capturing from webcam in OpenCV

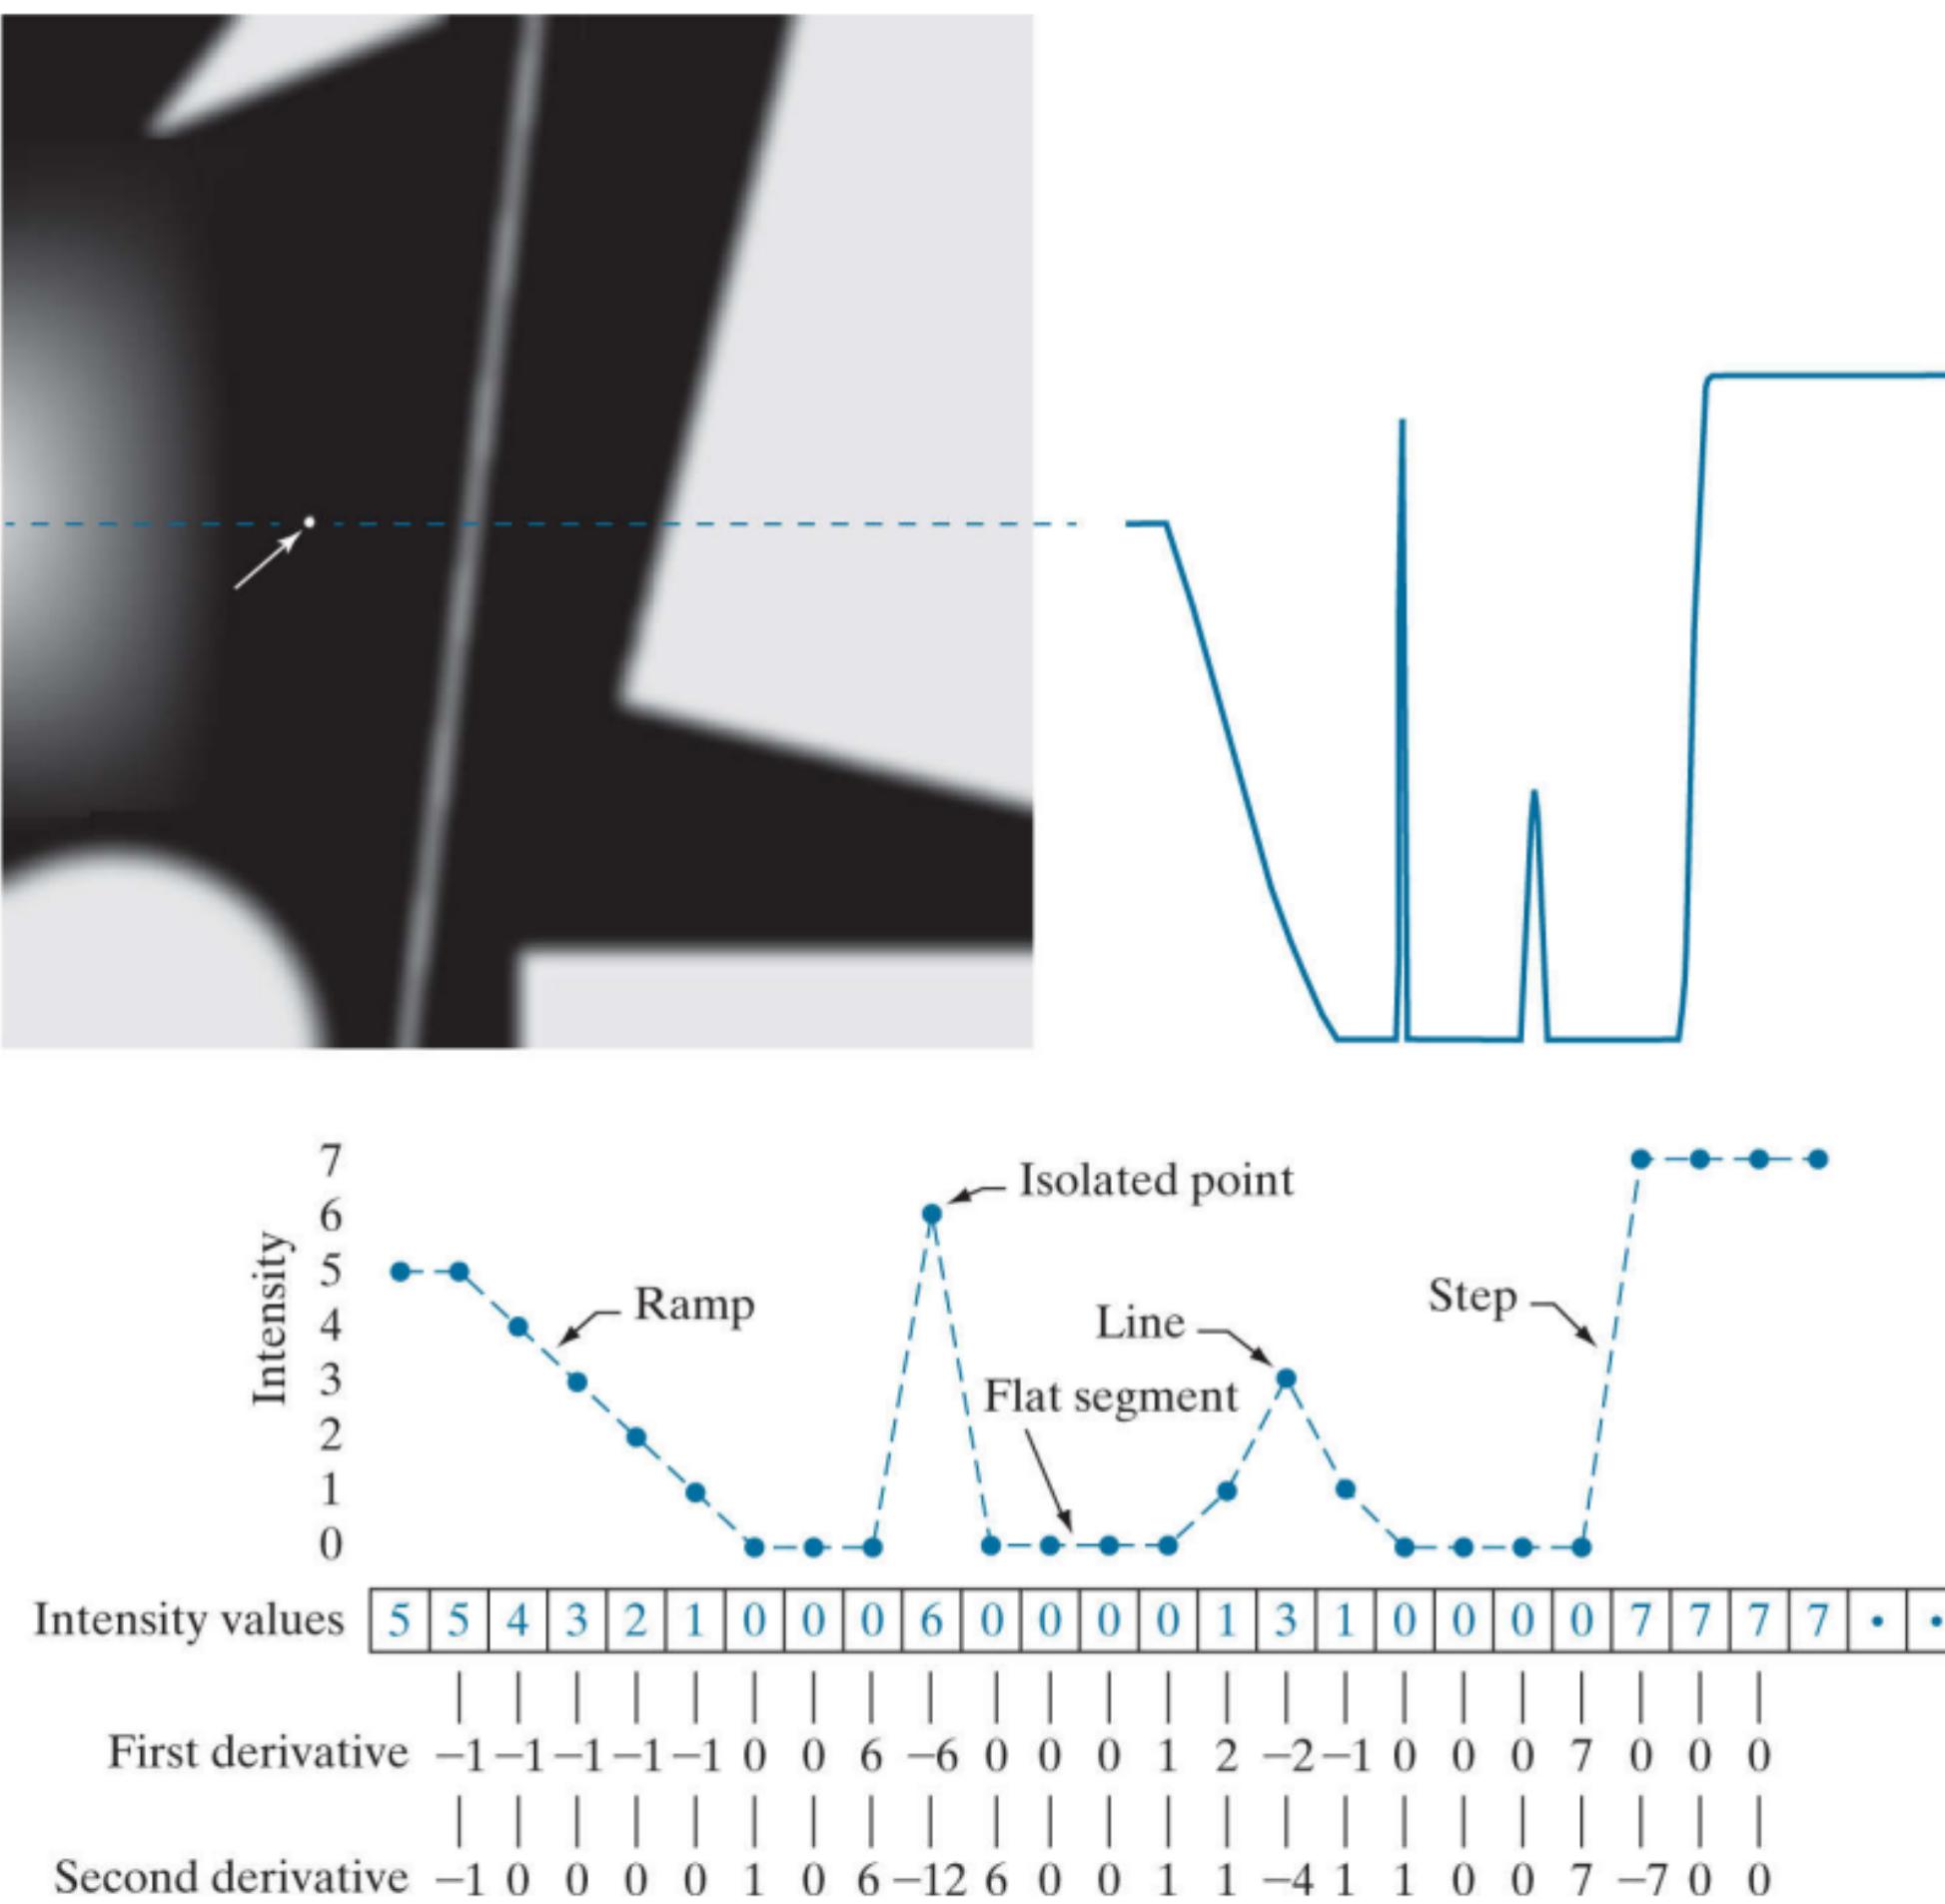
```
import cv2

cap = cv2.VideoCapture(0)
while True:
    return_val, frame = cap.read()
    cv2.imshow("Showing video", frame)
    if(cv2.waitKey(1) & 0xFF == ord('q')):
        break
```

0 - Open Default Camera

return\_val is True/False  
depending on success/failure  
of reading the video

# Image Segmentation



## First Order Derivative

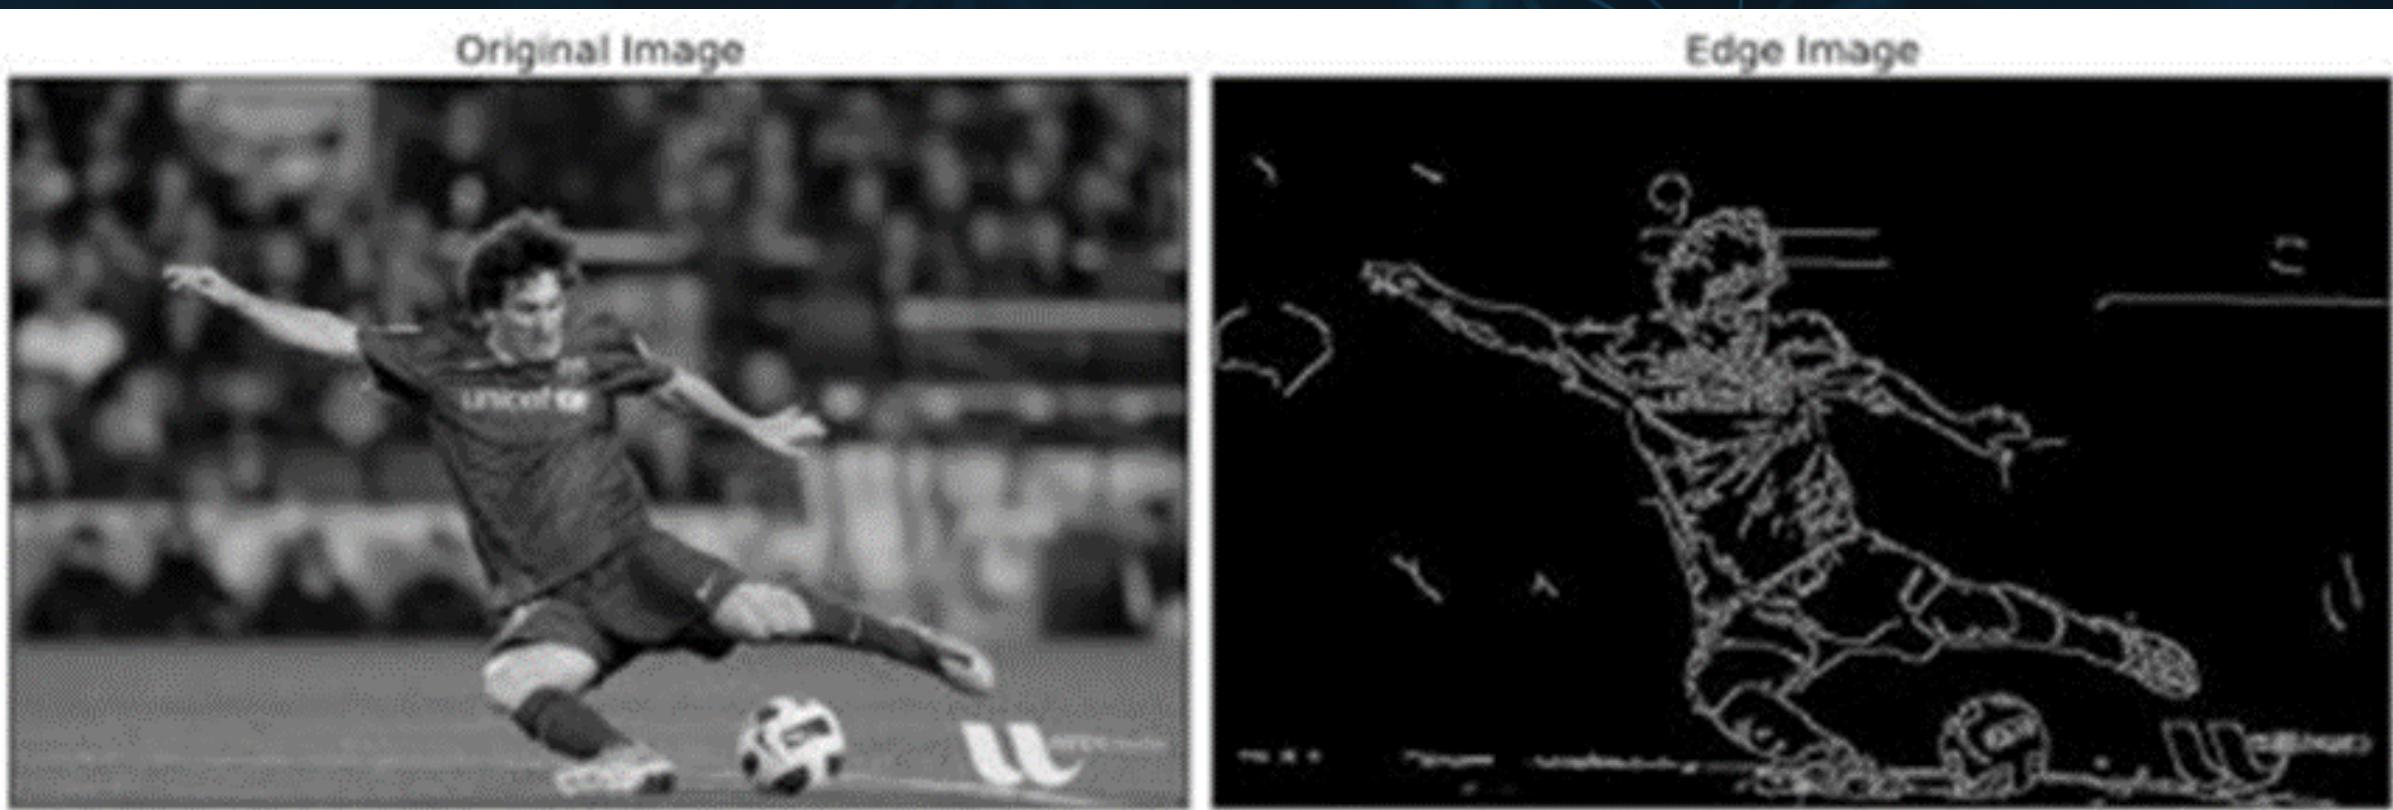
- thicker edges

## Second-Order Derivative

- fine details - isolated points, line
- double edge response at ramp and step
- to determine transition from light to dark or dark to light

# Edge Detection

- Canny Edge Detection
- Works on gray scale images



```
return_val, frame = cap.read()  
image = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)  
edges = cv2.Canny(image, 100, 200)
```

Maximum length

Minimum Length

# Contour Detection

A Contour is a closed curve joining all continuous points having same color or intensity. This is a important tool in shape analysis and object detection.



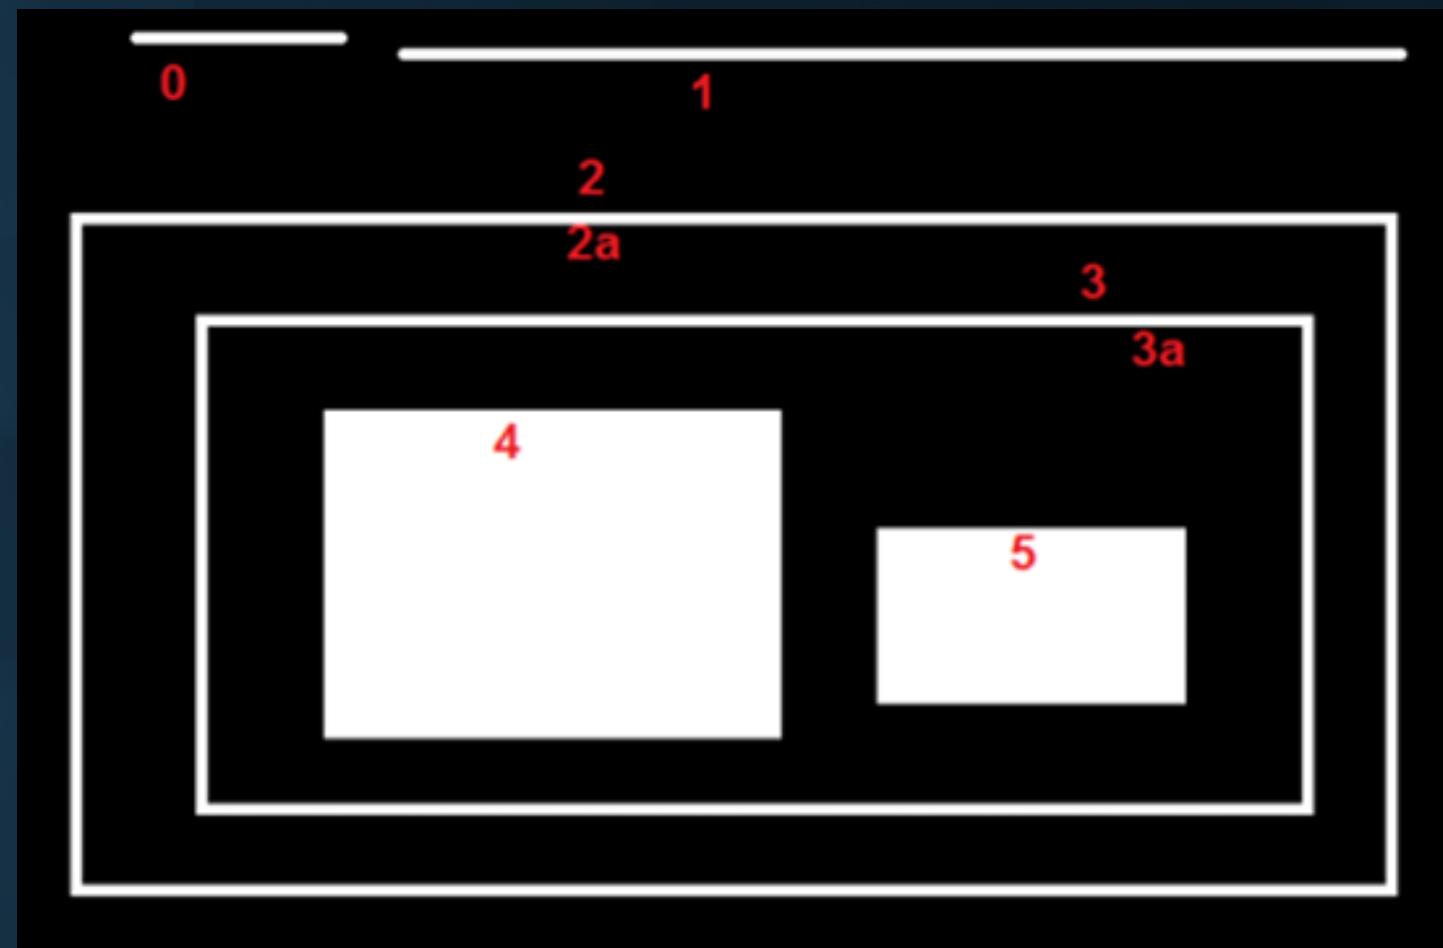
# Difference ?

Edge Detection and Contour Detection



# Contour Detection

Better accuracy is obtained when we use binary images. So before finding contours we can apply thresholding or canny edge detection.



`findContours(img, mode, method)`

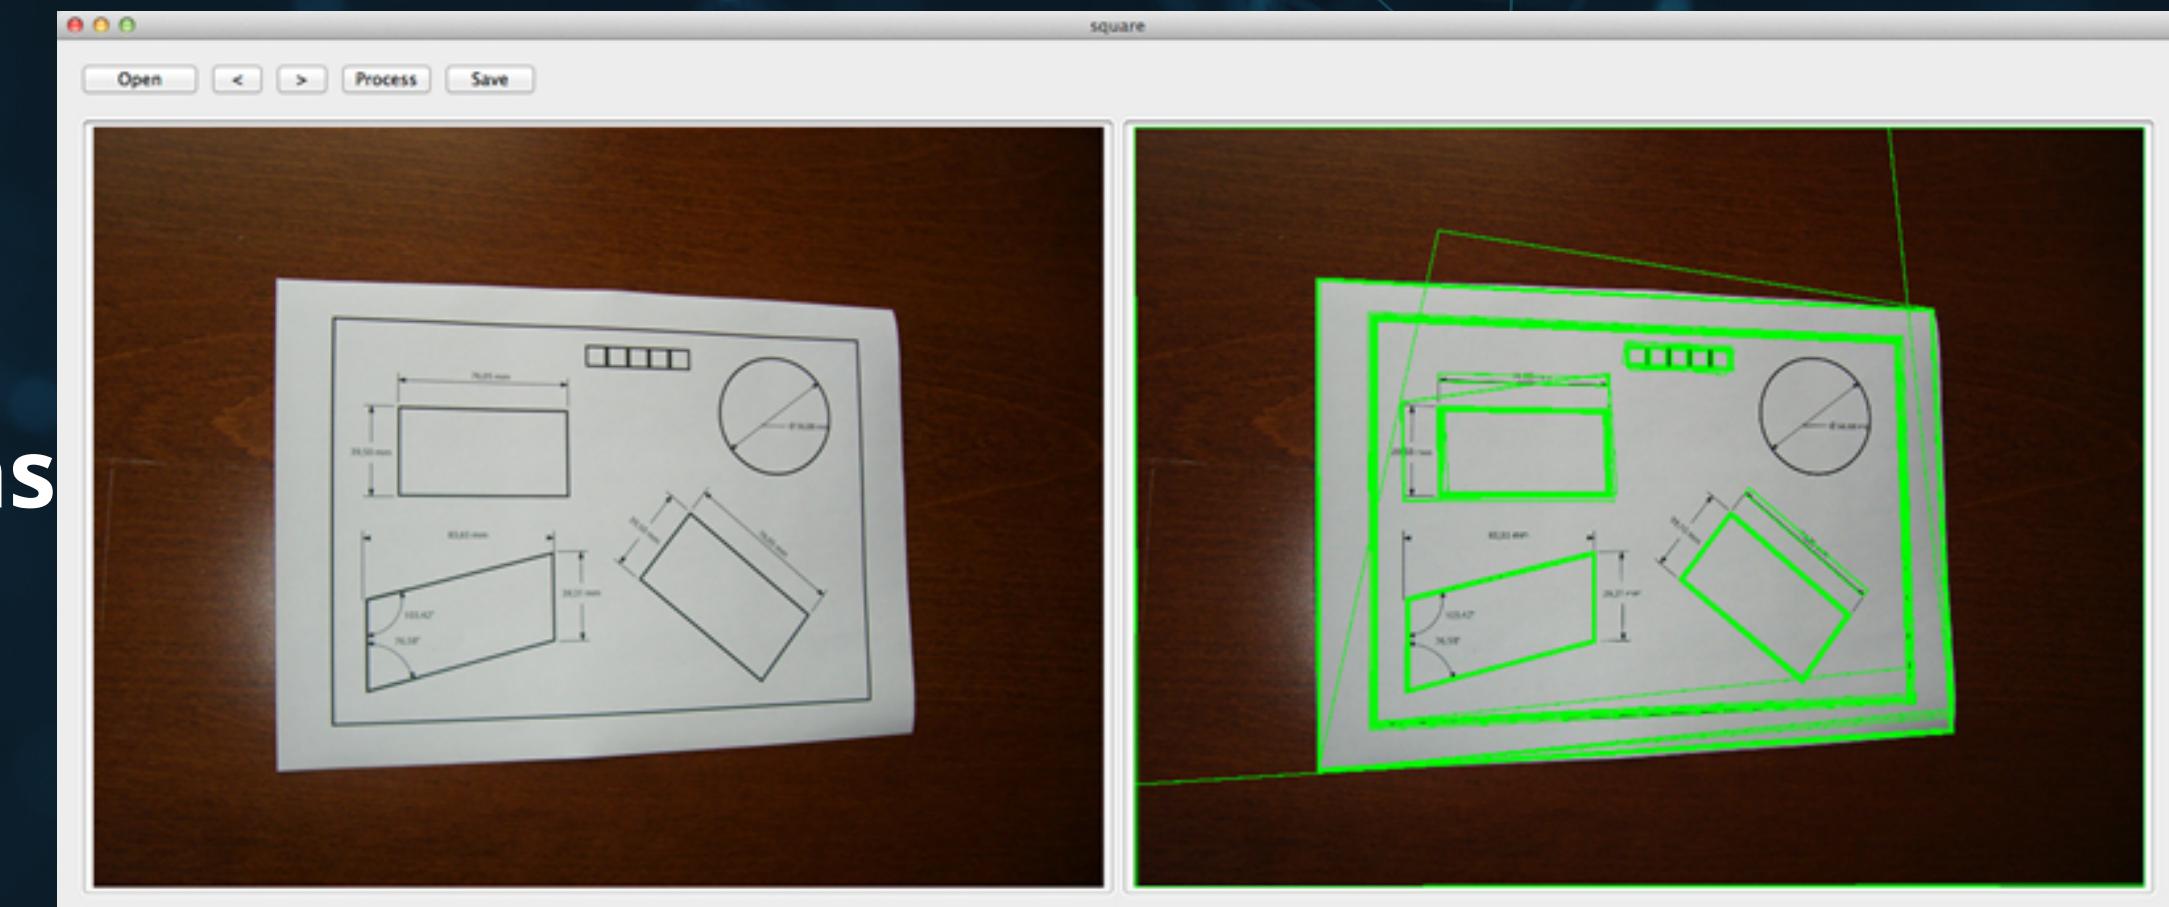
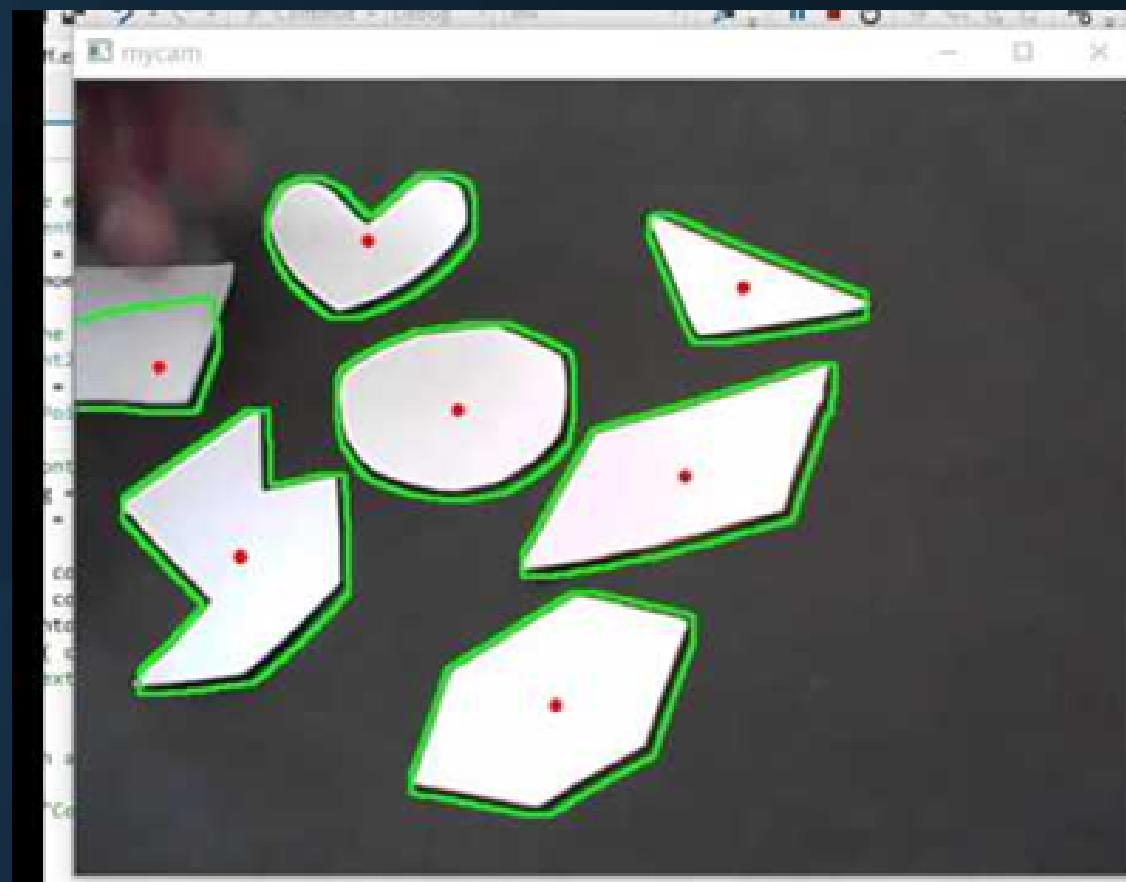
`img`: The image to be thresholded

`mode`: retrieval method (hierarchy)

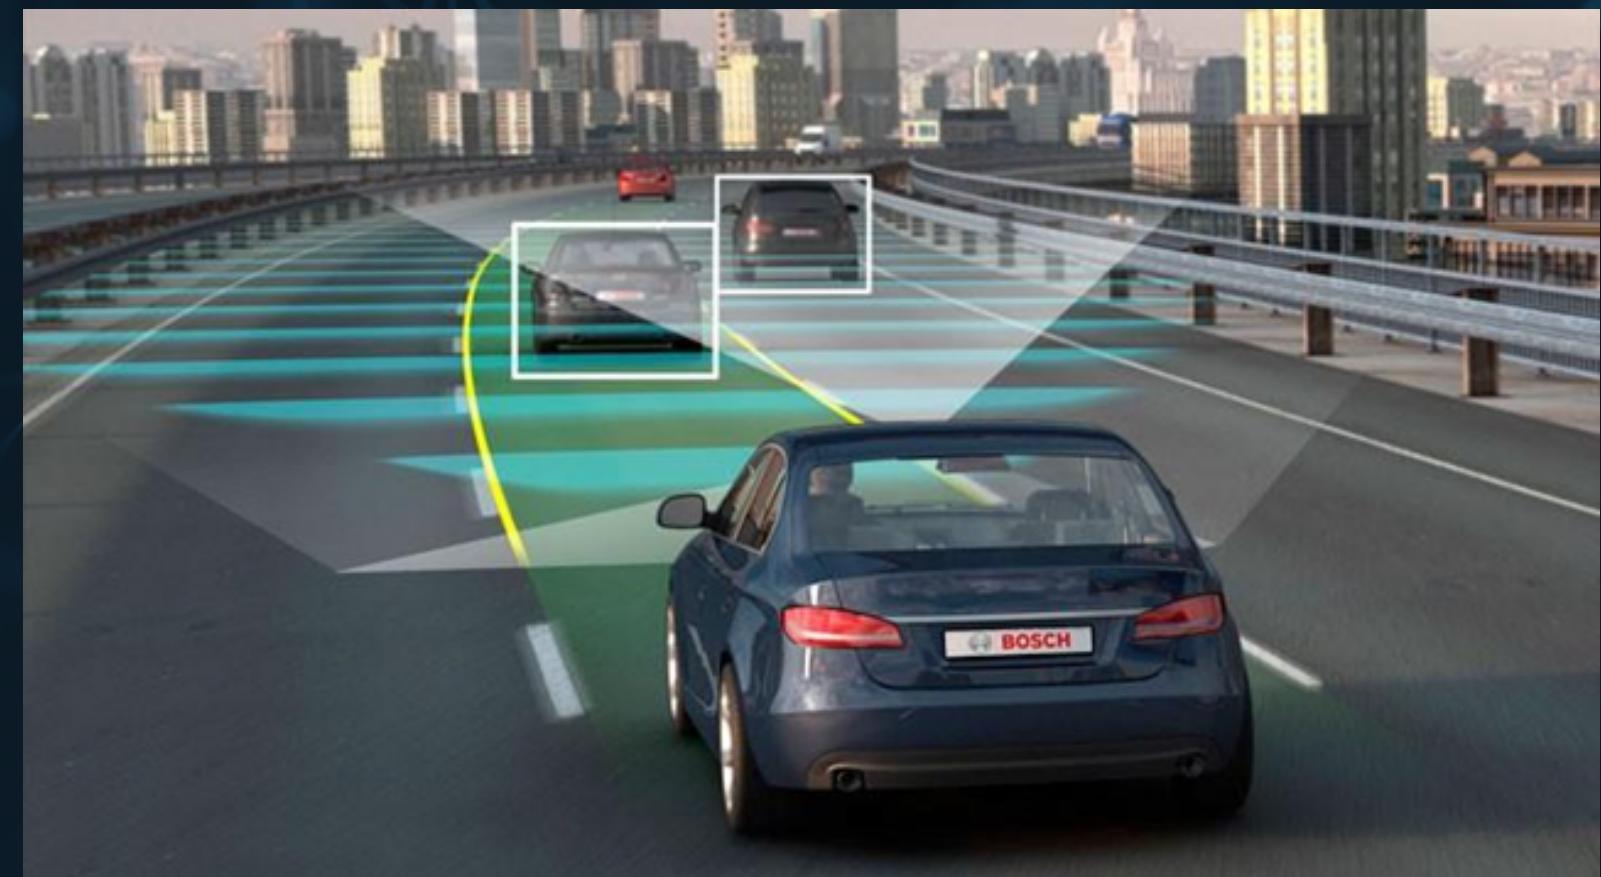
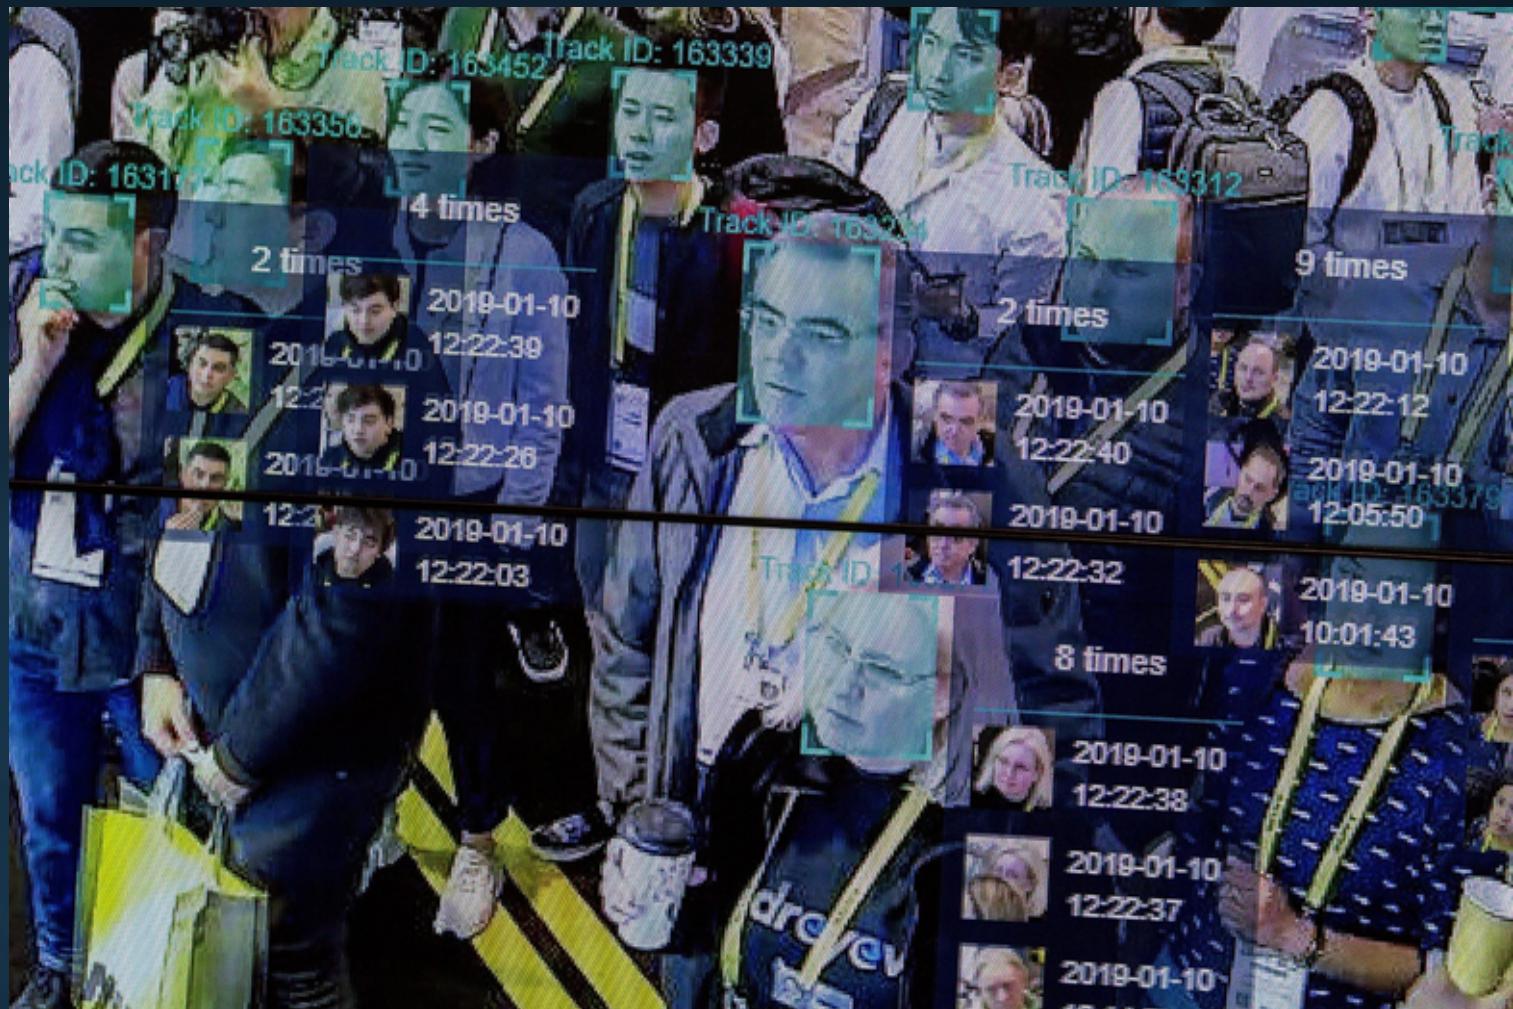
`method`: approximation method (corners)

# Contour Detection - Things to explore

- Perimeter
- Area
- Moments (to find centre)
- Approximation using polygons



# Applications of Image Processing/Computer Vision

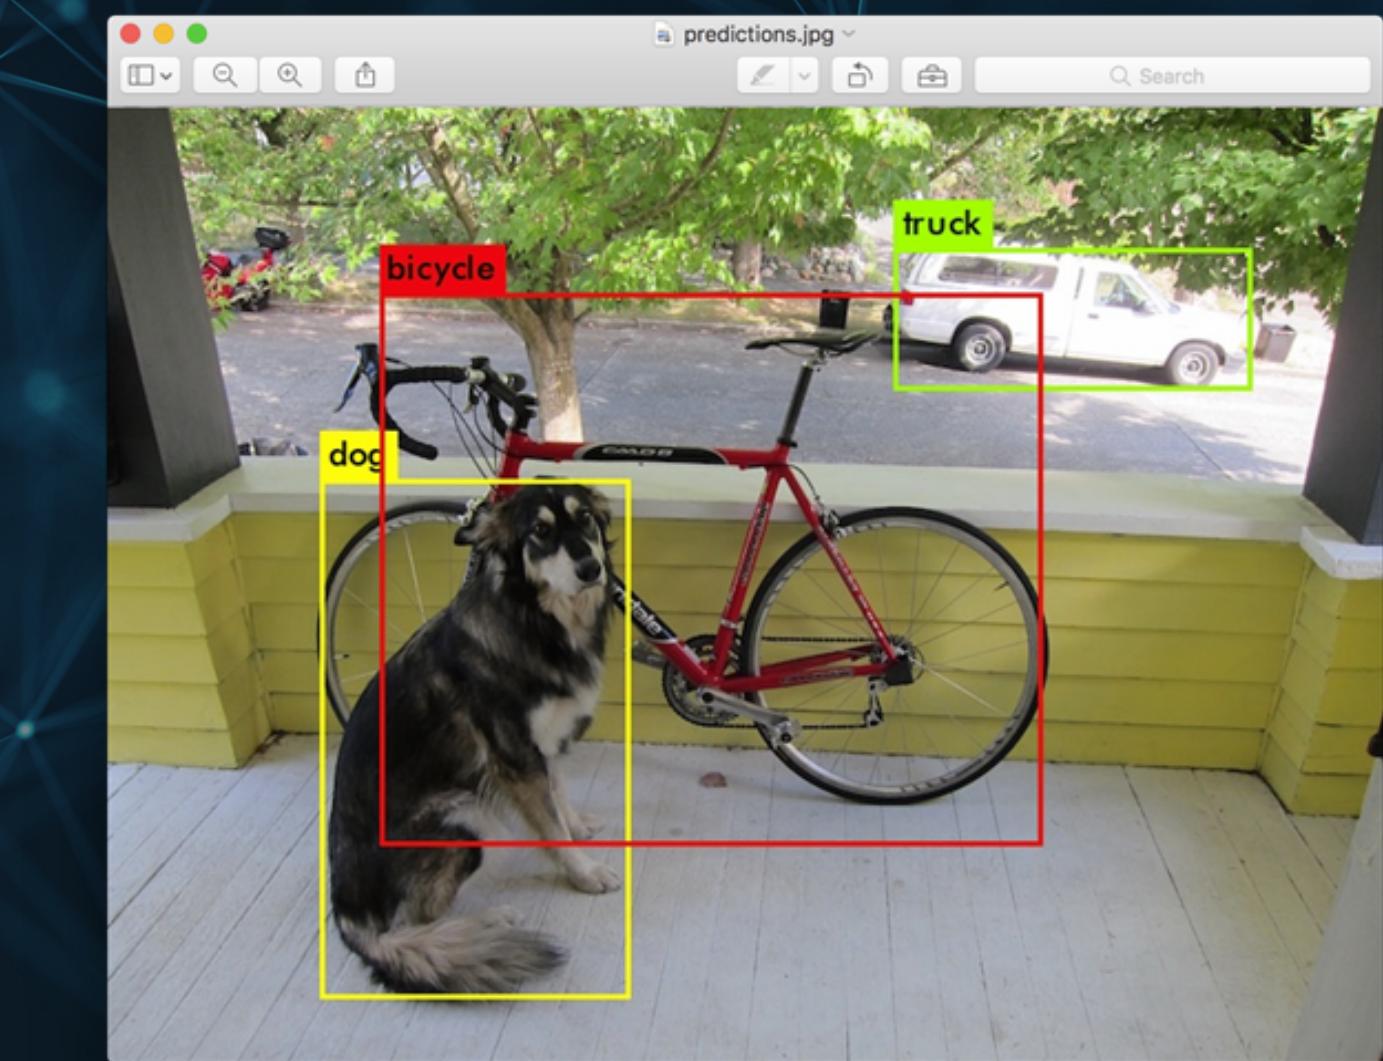


# Machine Learning and other advances

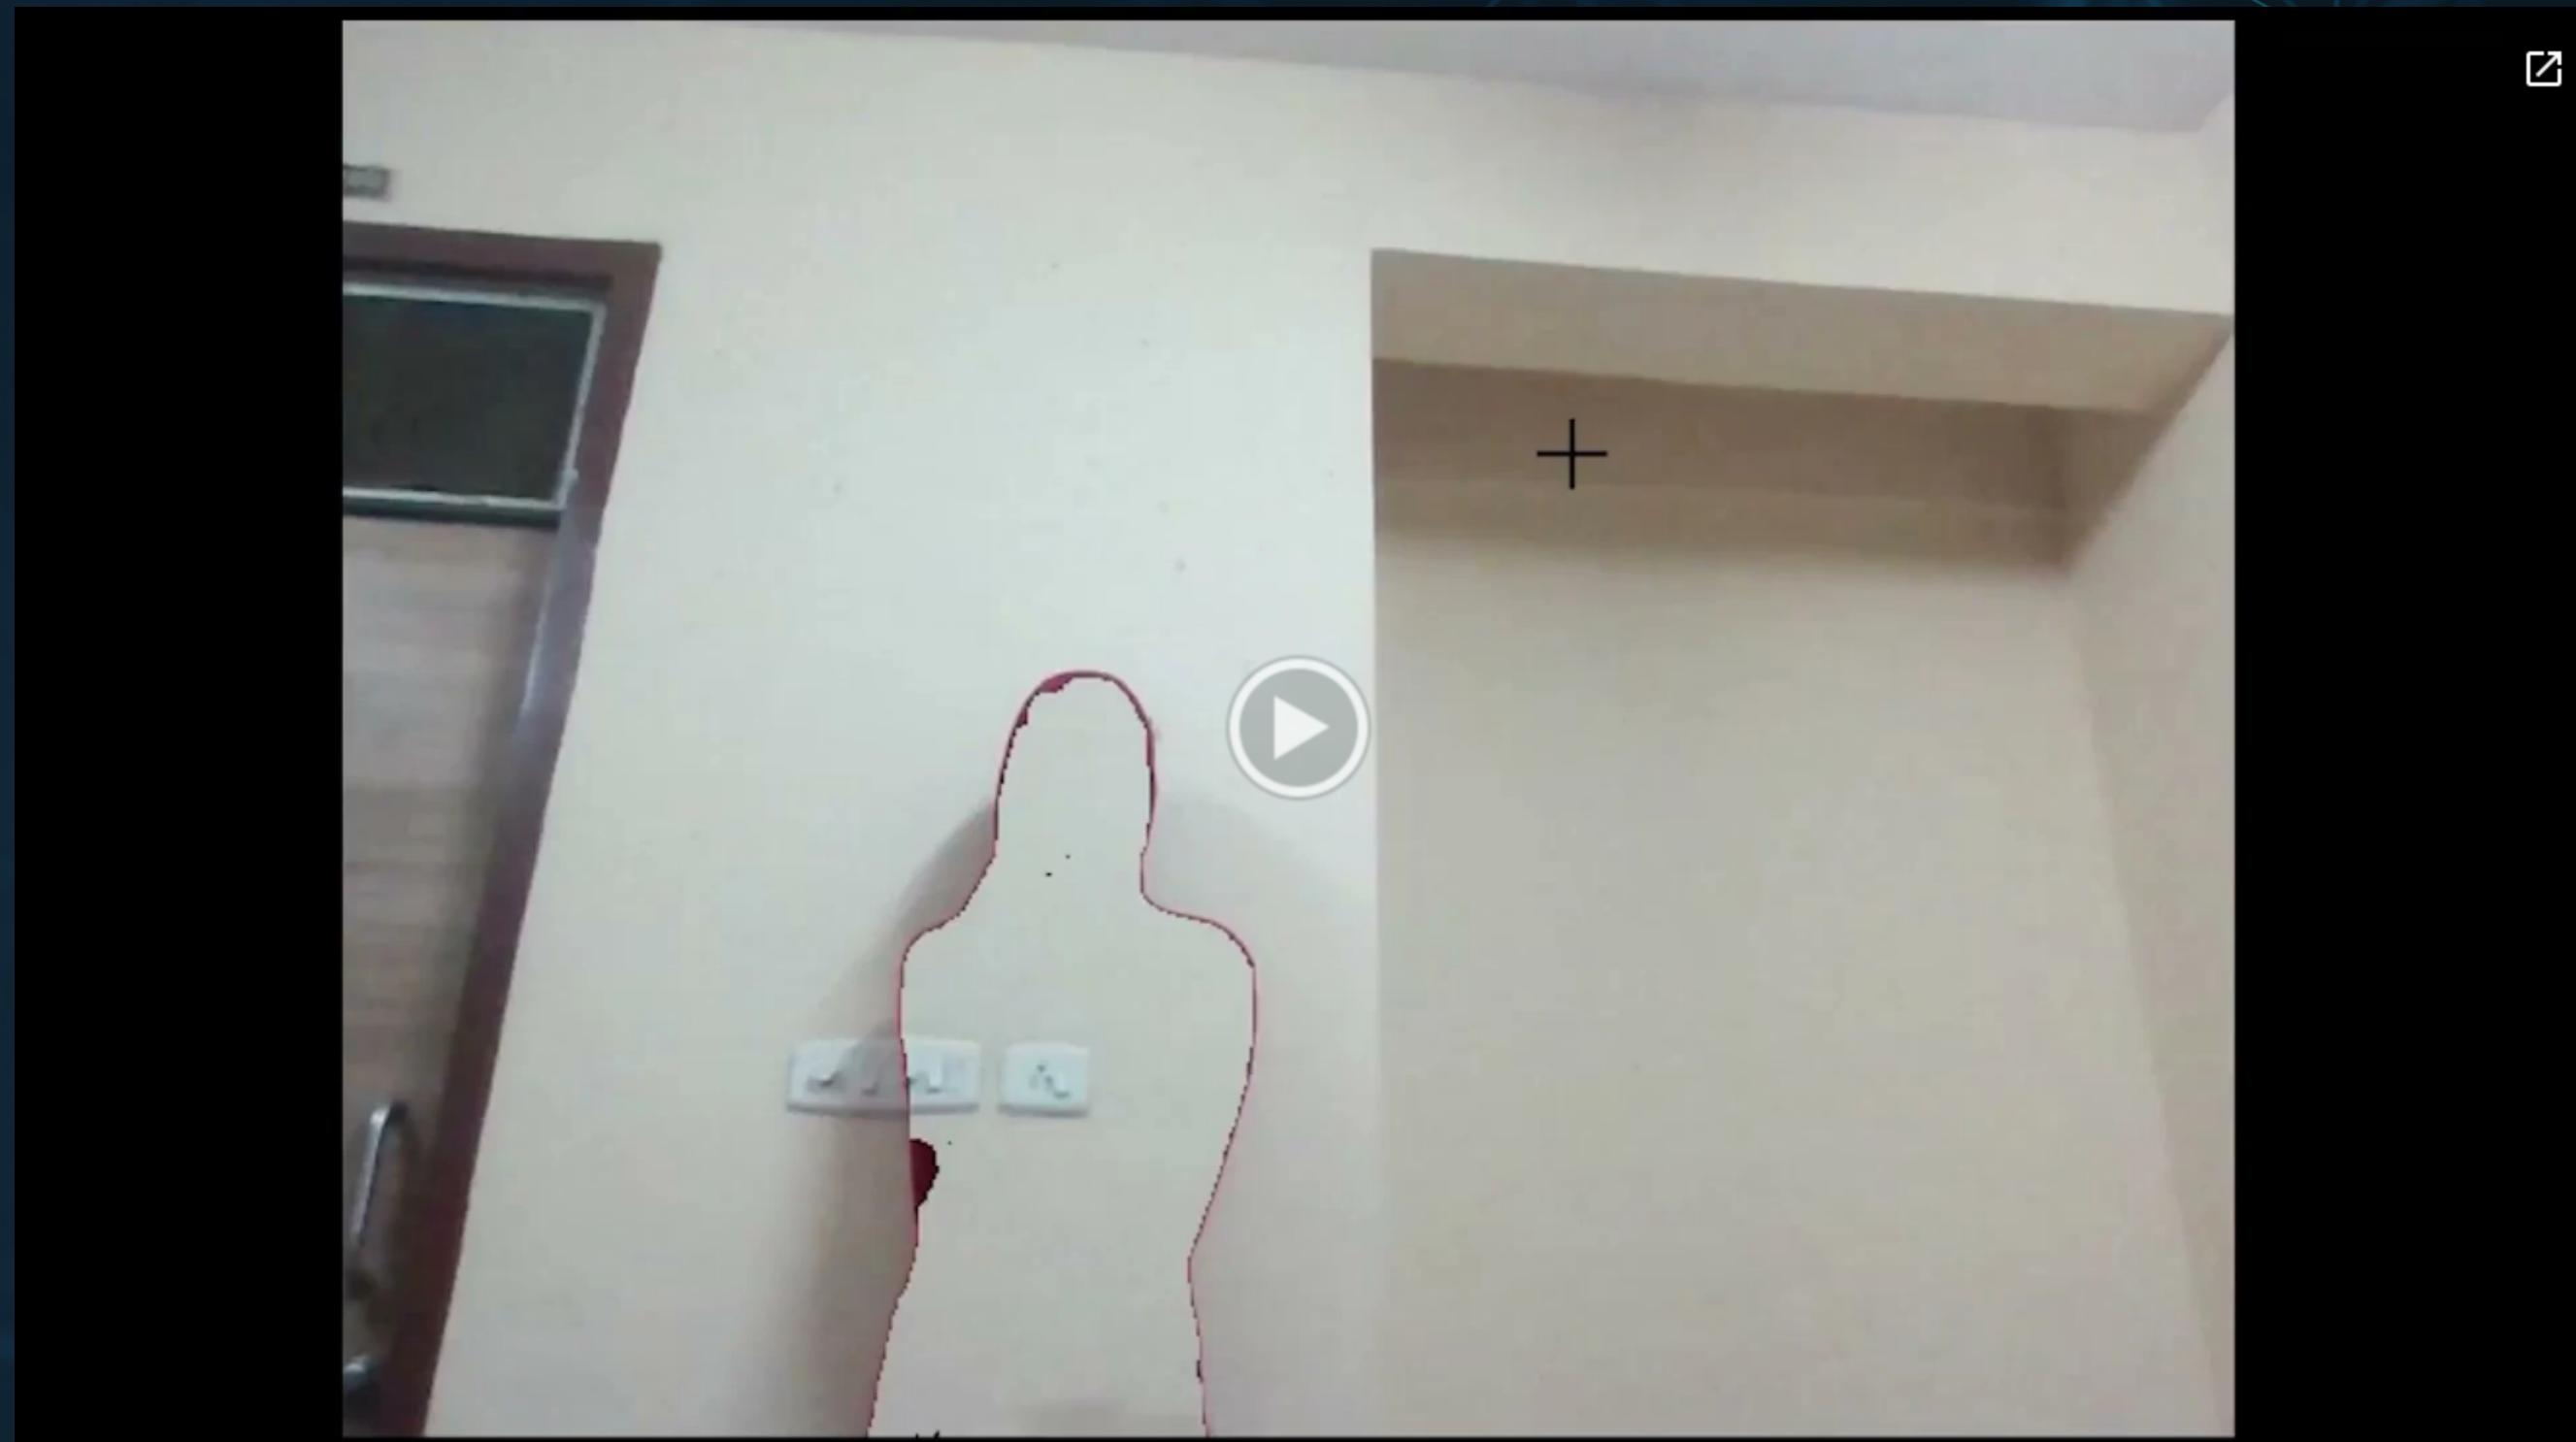
<https://pjreddie.com/darknet/yolo/>

<https://github.com/puzzledqs/BBox-Label-Tool>

<https://github.com/pytorch/vision/>



# Cool things to try out



# Thank You !